

TEALS | AI Summer Academy

Azure Global



Introduction

- **Azure Global Commercial Industry – RoW**
- **What we do** – Building industry PaaS AI solutions for the following:
 - Retail & Consumer Goods
 - Supply Chain
 - Energy
 - Financial Services
 - Media, Entertainment & Gaming

Python crash course

Emmanuel Awa

Said Bleik

Max Kaznady



Why Python

- **Learning Curve**
- **Awesome for POC and Scripting**
- **Platform Support and Flexibility**
- **A library/framework for everything**
- **Testing frameworks**
- **Great Choice for Enterprise Application Integration**
- **Active Technical Community**

Python Syntax

- Type into command line or file
- Indentation
- Variables
- Comments
- Casting
- Multiple assignments, variable unpacking
- Global vs Local Variables

Python Data Structures

- Numeric
 - Int, float, complex
- Sequence
 - List, tuple, range
- Mapping
 - Dict
- Sets
 - Set, frozenset
- Boolean
 - Bool
- Binary
 - Bytes, bytearray, memoryview
- Text
 - str

Python Operators

- Arithmetic
- Assignment
- Comparison
- Logical
- Identity
- Membership
- Bitwise

Conditionals and Loops

- Logical conditionals
 - If, if...elif..else
- Loops
 - While and For
 - Break statements are used to terminate loops

Functions and Lambdas Introduction

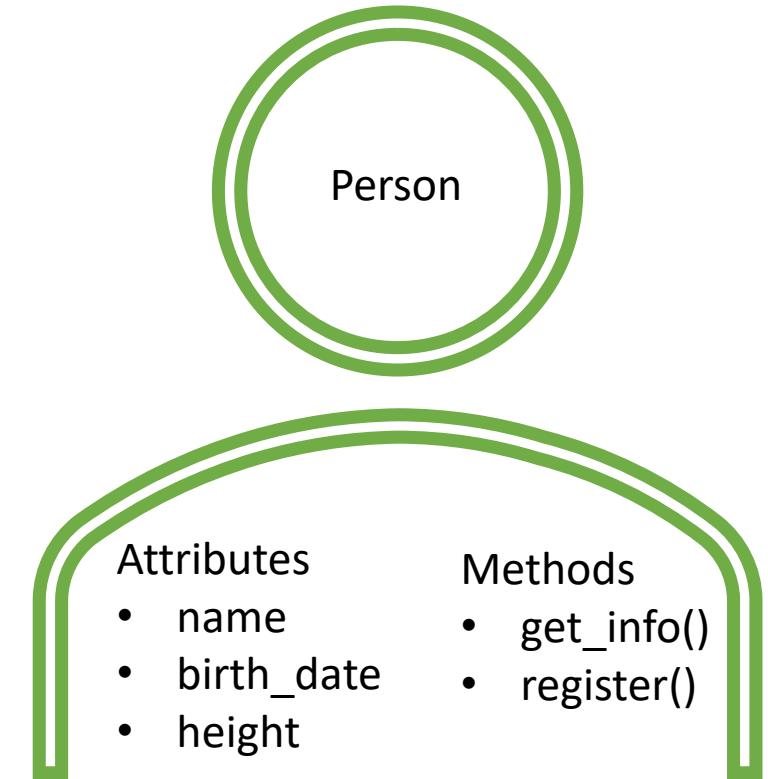
- Function is a named block of code that can only run when called.
 - It can take parameters by name or arbitrarily using *args
 - Can also take keyword arguments using **kwargs
- Lambda is a nameless function.
 - Can take many arguments but can only have one expression
 - Useful for quick reusable one-liners that are required for a short time
 - Faster than functions as it's evaluated at runtime

Go to notebook

Object Oriented Programming (OOP)

Beyond variables and values

- Programming paradigm where entities are represented as **objects**
- Abstract representation of properties and behavior of entities using **classes** (types of objects)
 - Properties <- member attributes
 - Behavior <- member methods
- Helps programmers:
 - Reuse code
 - Design better real-world systems
 - Build on top of others' work



Class **instances** (objects of type Person):
person1, person2, p3, that_person...

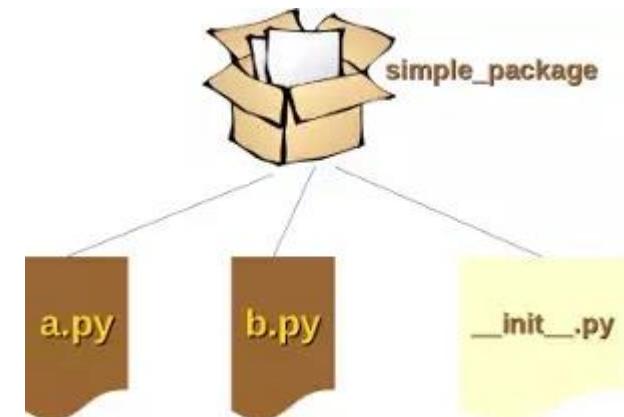
Object Oriented Programming (OOP)

- Inheritance
 - Parent/Child or Base/Derived classes
 - Ex: *class Student(Person)* and *class Teacher(Person)*
 - Inherit attributes and methods
- Encapsulation
 - Values and functions are encapsulated within a class
 - Ex: `student1.get_info()` returns previously set member attributes
- Polymorphism
 - Allows functions to have different forms
 - Ex: `student1.get_courses()` and `teacher5.get_courses()`

Go to notebook

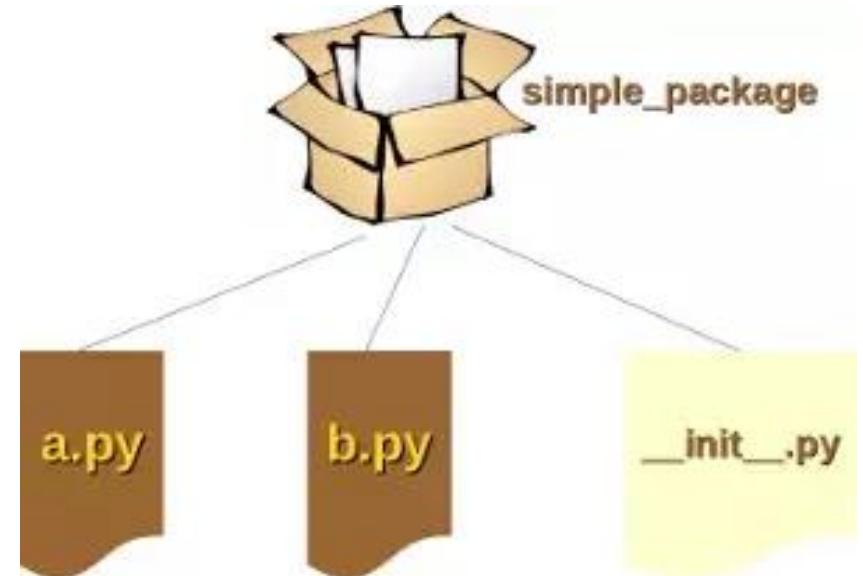
Packages

- Package is like a Toolbox
 - imports are like tools
- Toolbox example 1
 - “from toolbox import hammer
 - “hammer.strike()”
- Toolbox example 2
 - “from toolbox.hammer import strike”
 - “strike()”



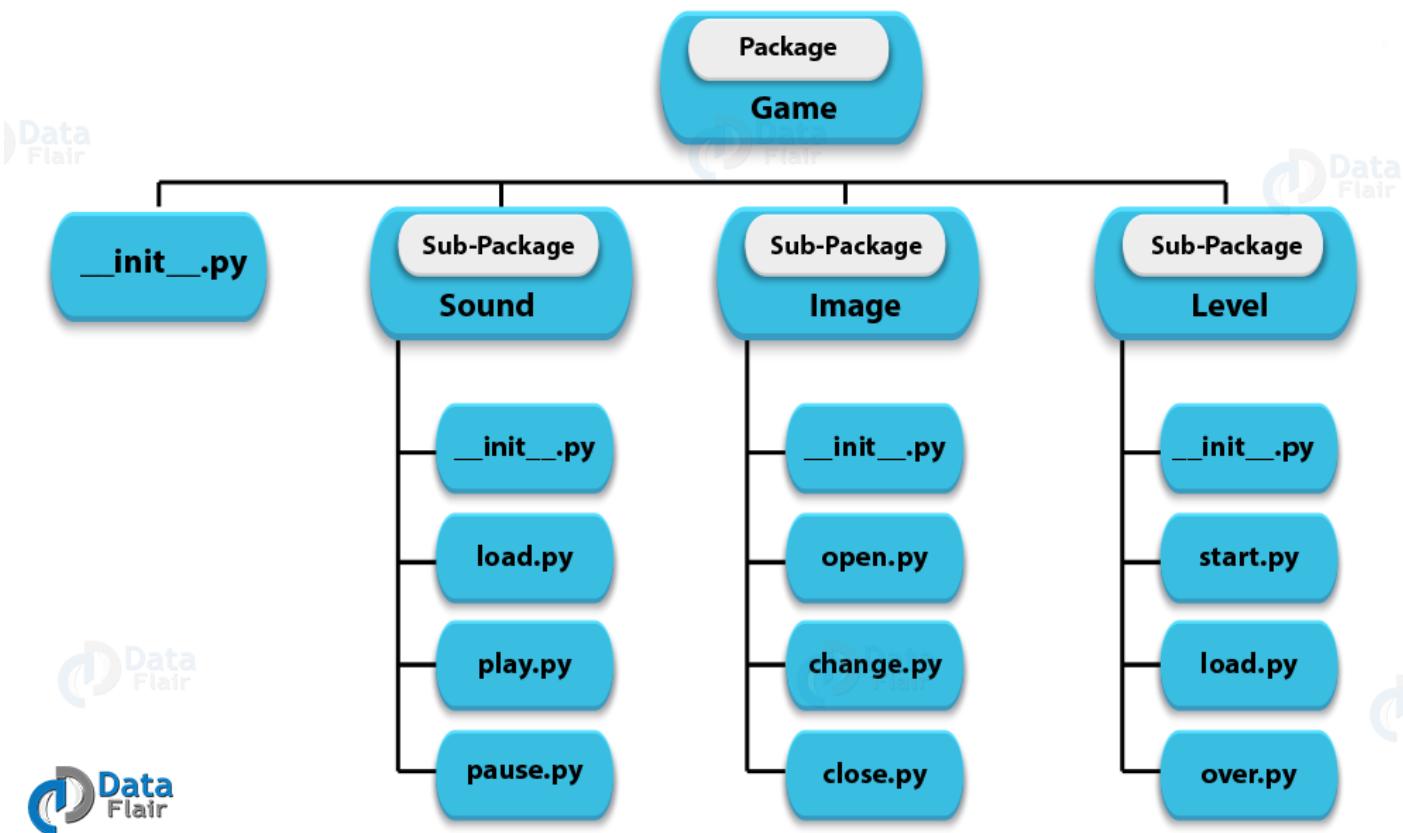
Package concepts

- Package module import
 - “from simple_package import a”
 - “a.do_stuff(‘some argument’)”
- Package function import
 - “from simple_package.a import do_stuff”
 - “do_stuff(‘some_argument’)”
- Package constants
 - “import simple_package; simple_package.CONSTANT”
 - “from simple_package import CONSTANT”



Growing packages

Package Module Structure



Coding example in an IDE

References

- [Free Intro to Python Course | Udacity | Udacity Free Courses](#)
- [Why Learn Python - Top 10 Reasons to Learn Python in 2021 | upGrad blog](#)
- [Python Tutorial \(w3schools.com\)](#)
- [\[https://www.python-course.eu/python3_packages.php\]\(https://www.python-course.eu/python3_packages.php\)](#)
- [<https://projectmanagementshop.eu/products/change-managers-toolbox/>](#)
- [<https://data-flair.training/blogs/python-packages/>](#)

AI Fundamentals

Alejandro Buendia

Jaya Mathew

Sahitya Mantravadi



Overview

Topics for the day

- Introduction to AI
- Probability and Statistics
- Linear Regression
- Binary Classification
- Q&A

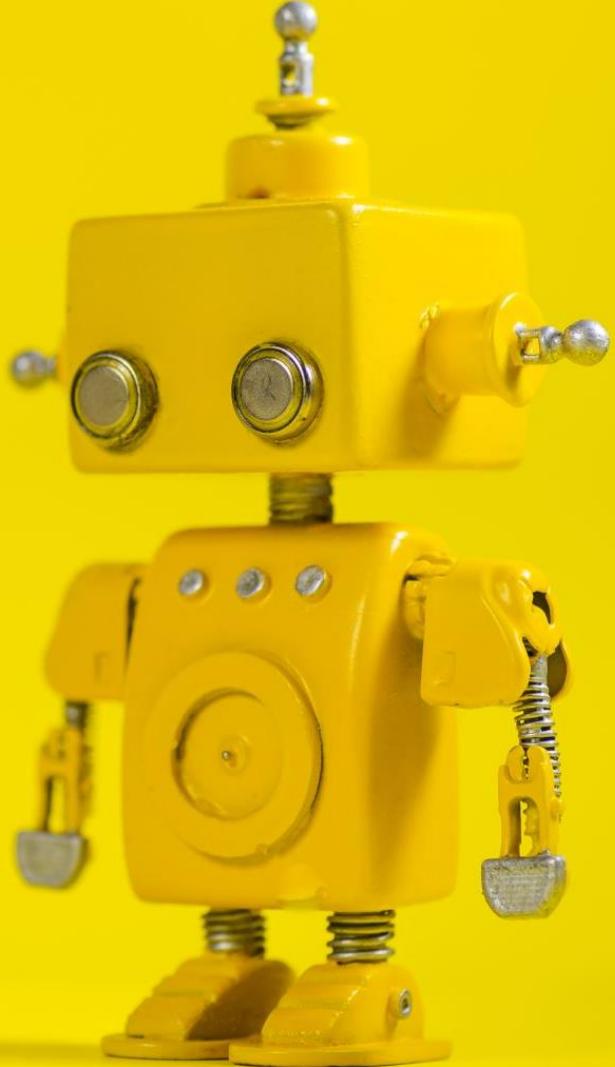
Speakers

- Alejandro Buendia – Introduction to AI, Probability and Statistics
- Jaya Mathew – Linear Regression
- Sahitya Mantravadi – Binary Classification

Introduction to AI

* Some graphics courtesy of Nakul Verma (COMS 4771, Columbia University)





What is AI and machine learning?

- Artificial intelligence (AI): Broadly refers to machines that can do tasks “intelligently.”
- Machine learning: Study of making machines learn a concept **from data** without having to explicitly program it.
- Machine learning can be thought of as a subset of AI.
- AI/machine learning involves creating algorithms that can:
 - Learn from a set of input data to make predictions
 - Find interesting patterns in data

Why AI/machine learning?

- Often we can't write down a simple set of rules to solve complicated real-world problems.
- For example, how would we write python code to distinguish between these two faces?



- Other times, we don't know what explicit task we want to solve given some data as a starting point.

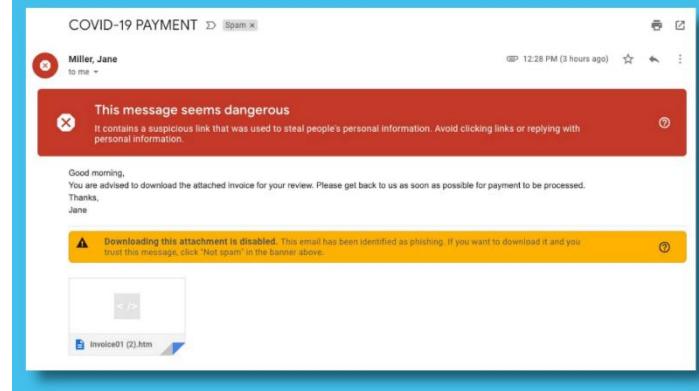
Machine Learning: The basics

Let's look at three prediction problems:

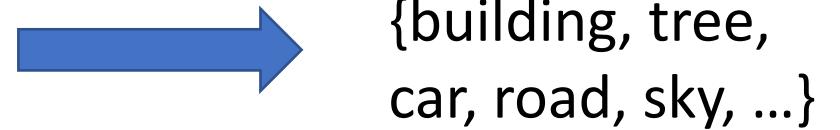
- Handwritten character recognition



- Spam filtering

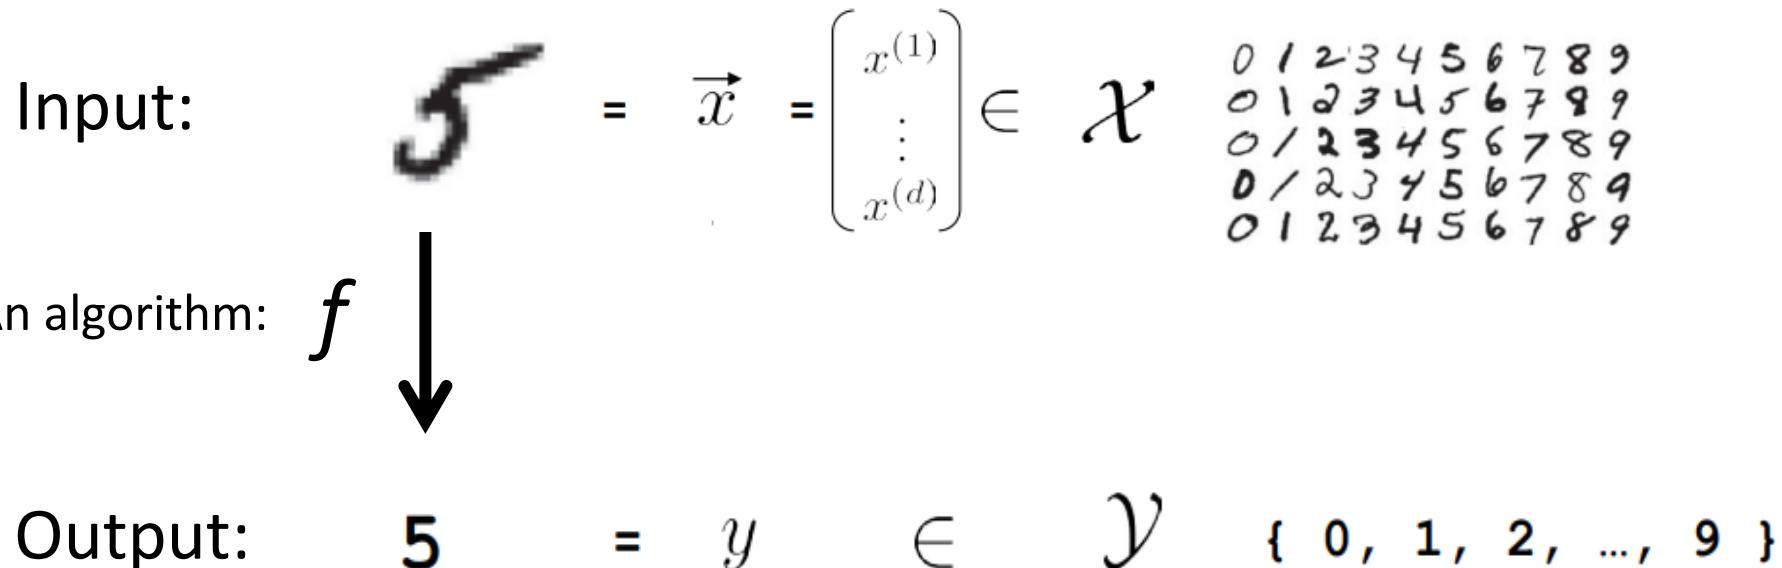


- Object recognition



Machine Learning: The basics (cont'd)

What do these problems have in common? We have...

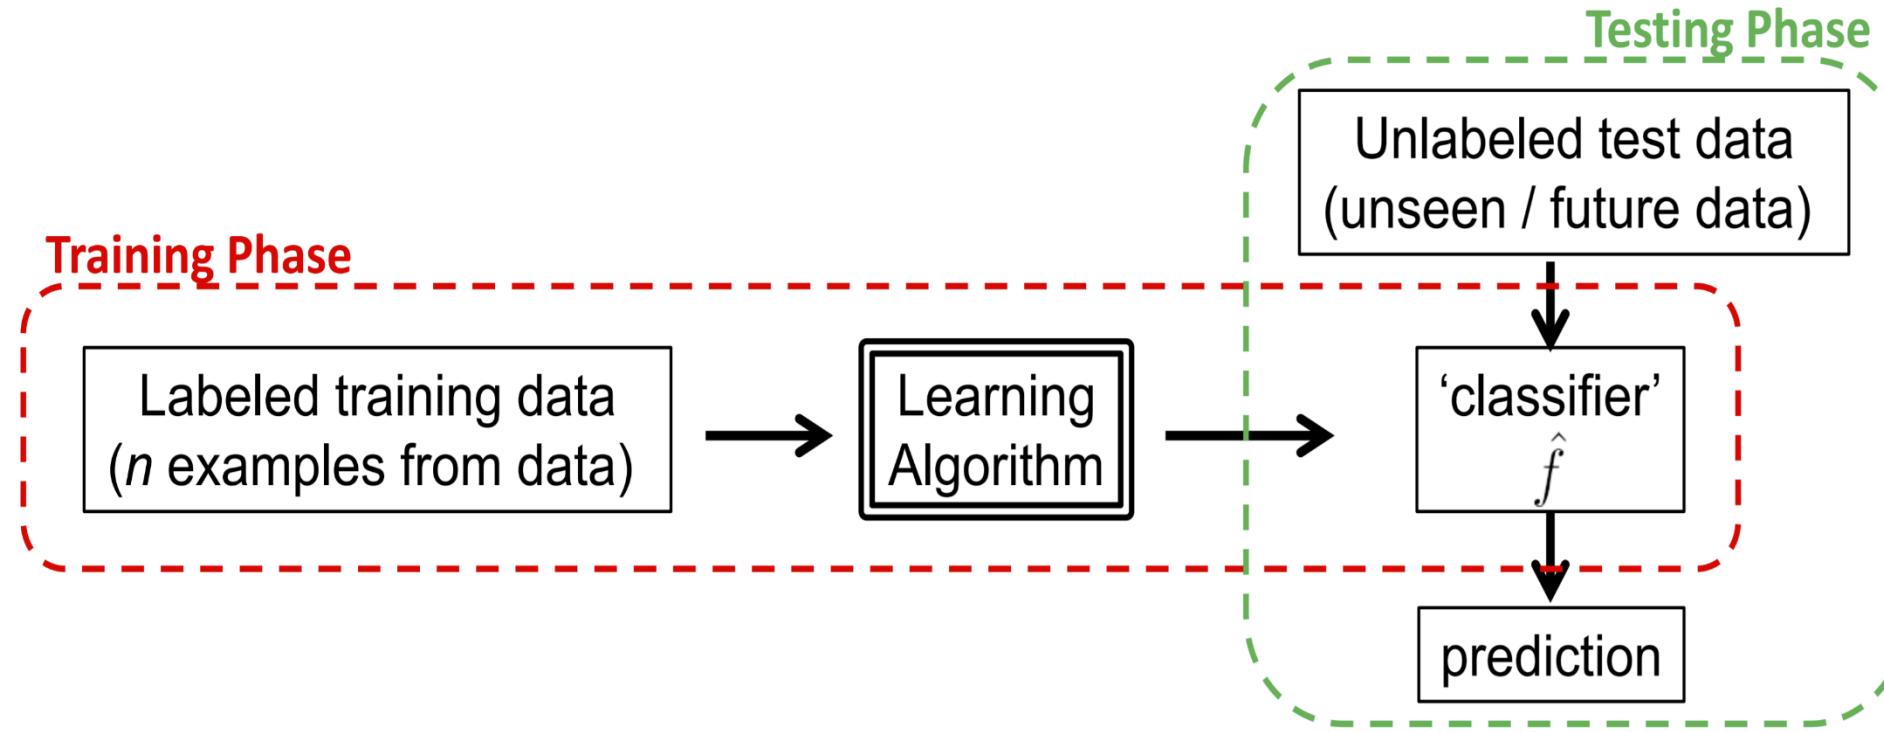


We want the machine to **learn** some function (algorithm)
 $f: X \rightarrow Y$ which maps the input set to the output set.

Machine Learning: The basics (cont'd)

- In **supervised** machine learning, we have a bunch of input and output pairs that we use to **train** the machine: $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
- Assumption: There is a function (algorithm) $f: X \rightarrow Y$ such that $f(x_1) = y_1, f(x_2) = y_2$, and so on, for **all** of the data points.
- f is the “real” answer (algorithm) that we want the machine to learn.
- Using n data points, we want to find an approximation \hat{f} for the “real” answer f .
- Goal: The approximation \hat{f} that the machine learns gives mostly correct predictions on unseen data.

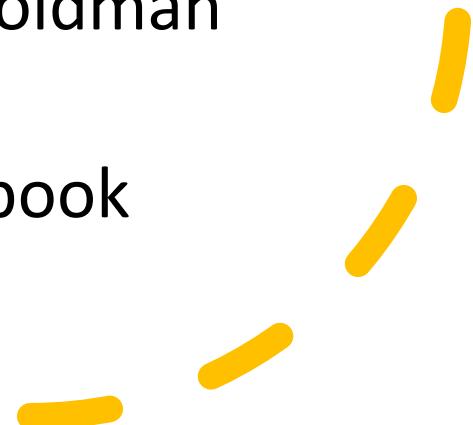
The machine learning workflow



AI problems in the real world

AI is everywhere!

- Search engines and content filtering: Google, Bing
- Product recommendation systems: Netflix, Amazon
- Speech recognition: Alexa, Siri
- Finance and predicting stocks: Goldman Sachs, Morgan Stanley
- Face or object recognition: Facebook





Now that we understand the problem setup, *how* do we actually have the machine learn the model \hat{f} ?

We'll look at a few different approaches.

Probability and Statistics





What is a random variable?

- A **random variable** X is a variable whose values are outcomes of a random phenomenon.
- **Discrete** random variables take on a countable number of distinct values. For example, the value of a coin flip $X = \{\text{heads, tails}\}$.
- **Continuous** random variables take on one of an infinite number of possible values. For example, X could represent height and take on any positive value: 5.5 feet, 5.7 feet, etc.

How do we compute probabilities?

- The **probability** that a random variable X takes on a particular value A describes how likely X will be equal to A , written as $P(X = A)$.
- Probability is a value between 0 and 1, where 0 means the event never happens and 1 means the event always happens.
- For example, $P(X = \text{heads}) = \frac{1}{2}$ and $P(X = \text{tails}) = \frac{1}{2}$ if X describes a regular, fair coin. We compute this by taking the outcome of interest over all possible outcomes:

$$P(X = \text{heads}) = \frac{|\text{heads}|}{|\{\text{heads, tails}\}|} = \frac{1}{2}$$

Conditional probabilities

- We can also look at **conditional probability**, which is the probability of one event A happening given that another event B happened, or $P(A | B)$.
- Often $P(A | B)$ is different from $P(A)$. For example, $P(\text{snow today} | \text{winter})$ is likely much different than $P(\text{snow today})$.
- We compute conditional probability as:

$$P(A | B) = \frac{P(\text{A and B})}{P(B)}$$

Joint probability
that both A and B
happen

Probability that B
happens

Bayes' rule

1. Using the previous definitions for conditional probabilities:

$$P(A | B) = \frac{P(A \text{ and } B)}{P(B)}$$

$$P(B | A) = \frac{P(A \text{ and } B)}{P(A)}$$

2. Rewrite as:

$$P(A \text{ and } B) = P(A | B) P(B)$$

$$P(A \text{ and } B) = P(B | A) P(A)$$

3. Set right hand sides equal to each other:

$$P(A | B) P(B) = P(B | A) P(A)$$

4. Divide by $P(B)$:

$$P(A | B) = \frac{P(B | A) P(A)}{P(B)}$$

← Bayes' rule

Bayes' rule in action...

Suppose that a particular cancer test is:

1. 97% correct for catching positive cases (3% of time gives a negative result when it should have been positive)
2. 95% correct for catching negative cases (5% of time gives a positive result when it should have been negative)

0.5% of the general population has this type of cancer. If a randomly sampled person tests positive (+), what is the probability that they actually have cancer?

$$\begin{aligned} P(\text{cancer} | +) &= \frac{P(+ | \text{cancer}) P(\text{cancer})}{P(+ | \text{cancer}) P(\text{cancer})} \\ &= \frac{P(+ | \text{cancer}) P(\text{cancer}) + P(+ | \text{no cancer}) P(\text{no cancer})}{0.97 * 0.005} \\ &= \frac{0.97 * 0.005 + 0.05 * (1 - 0.005)}{0.97 * 0.005 + 0.05 * (1 - 0.005)} = 0.089 = 8.9\% \end{aligned}$$

Why is Bayes' rule important?

We see a surprising result in the previous example: there's only a 8.9% chance that someone who tests positive actually has cancer!

This is because we incorporated the **prior** information that the prevalence of cancer is very low in the population, leading to a low probability.

Our first AI model: Naïve Bayes classifier

- We can apply Bayes' theorem to create an AI algorithm that predicts whether an incoming email is spam or not!
- Assume we have a dataset of 8 “not spam” emails and 4 “spam” emails which we found in our inbox.
- We went through and counted the number of times the following words appeared in each email type overall: {dear, friend, money}

	Not Spam (8 emails)	Spam (4 emails)
Dear	8	2
Friend	5	1
Money	1	4

Naïve Bayes classifier (Example 1)

- We receive a new email that reads “dear friend.” We want to determine whether this email is spam or not spam using an AI classifier.
- Let’s say that we are trying to predict variable y , which can be “spam” or “not spam.”
- We are interested in these two conditional probabilities:

$$\begin{aligned} P(y = \text{spam} \mid \text{email} = \text{"dear friend"}) \\ P(y = \text{not spam} \mid \text{email} = \text{"dear friend"}) \end{aligned}$$

- Whichever probability is higher will tell us our prediction for the email!

Naïve Bayes classifier (Example 1)

We can use Bayes' rule!

$$\begin{aligned} P(\text{spam} \mid \text{"dear friend"}) &= \frac{P(\text{"dear friend"} \mid \text{spam}) P(\text{spam})}{P(\text{"dear friend"})} \\ &= \frac{P(\text{"dear"} \mid \text{spam}) P(\text{"friend"} \mid \text{spam}) P(\text{spam})}{P(\text{"dear friend"})} = \frac{\frac{2}{7} * \frac{1}{7} * \frac{4}{12}}{P(\text{"dear friend"})} \propto 0.014 \end{aligned}$$

$$\begin{aligned} P(\text{not spam} \mid \text{"dear friend"}) &= \frac{P(\text{"dear friend"} \mid \text{not spam}) P(\text{not spam})}{P(\text{"dear friend"})} \\ &= \frac{P(\text{"dear"} \mid \text{not spam}) P(\text{"friend"} \mid \text{not spam}) P(\text{not spam})}{P(\text{"dear friend"})} = \frac{\frac{8}{14} * \frac{5}{14} * \frac{8}{12}}{P(\text{"dear friend"})} \propto 0.136 \end{aligned}$$

	Not Spam (8 emails)	Spam (4 emails)
Dear	8	2
Friend	5	1
Money	1	4

We conclude that the email is **not spam** since:

$$\begin{aligned} P(\text{not spam} \mid \text{"dear friend"}) &> P(\text{spam} \mid \text{"dear friend"}) \end{aligned}$$

Naïve Bayes classifier (Example 2)

Now the text in the email reads “money friend”:

$$\begin{aligned} P(\text{spam} \mid \text{"money friend"}) &= \frac{P(\text{"money friend"} \mid \text{spam}) P(\text{spam})}{P(\text{"money friend"})} \\ &= \frac{P(\text{"money"} \mid \text{spam}) P(\text{"friend"} \mid \text{spam}) P(\text{spam})}{P(\text{"money friend"})} = \frac{\frac{4}{7} * \frac{1}{7} * \frac{4}{12}}{P(\text{"money friend"})} \propto 0.027 \end{aligned}$$

$$\begin{aligned} P(\text{not spam} \mid \text{"money friend"}) &= \frac{P(\text{"money friend"} \mid \text{not spam}) P(\text{not spam})}{P(\text{"money friend"})} \\ &= \frac{P(\text{"money"} \mid \text{not spam}) P(\text{"friend"} \mid \text{not spam}) P(\text{not spam})}{P(\text{"money friend"})} = \frac{\frac{1}{14} * \frac{5}{14} * \frac{8}{12}}{P(\text{"money friend"})} \propto 0.017 \end{aligned}$$

	Not Spam (8 emails)	Spam (4 emails)
Dear	8	2
Friend	5	1
Money	1	4

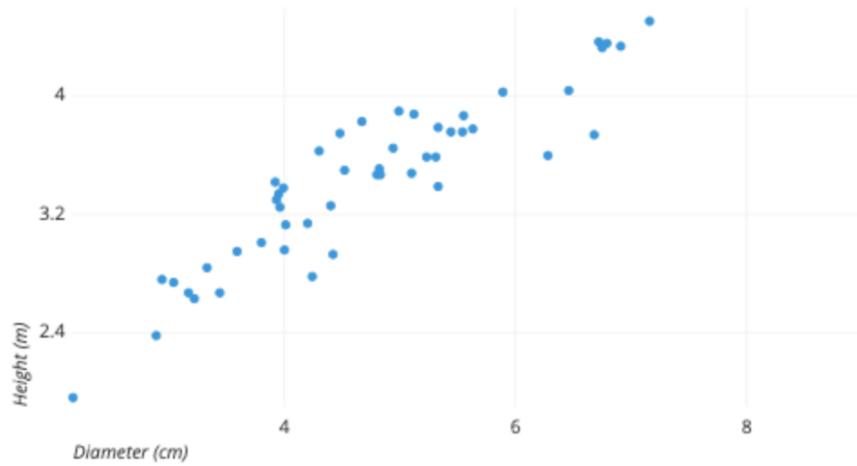
We conclude that the email is **spam** since:
 $P(\text{spam} \mid \text{"money friend"}) > P(\text{not spam} \mid \text{"money friend"})$



Linear Regression

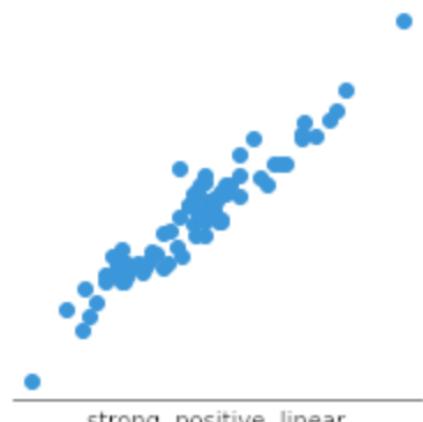
What is a scatter plot?

- A scatter plot (aka scatter chart, scatter graph) uses dots to represent values for two different numeric variables.
- The position of each dot on the horizontal and vertical axis indicates values for an individual data point.

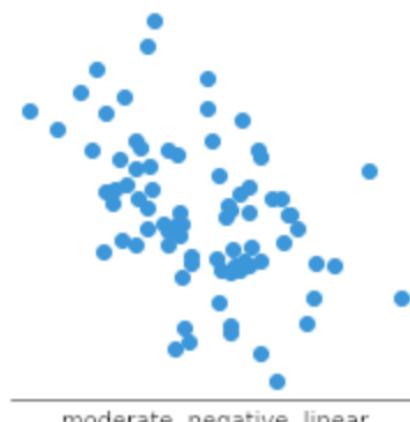


Let's look at some scatter plots

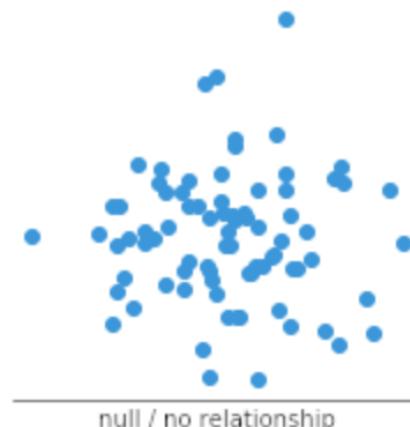
- Scatter plots are used to observe relationships between variables.
- In the plots, as you can see some of them seem to show some relationship between the variables



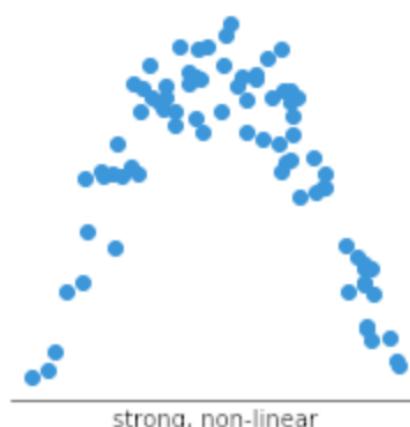
strong, positive, linear



moderate, negative, linear



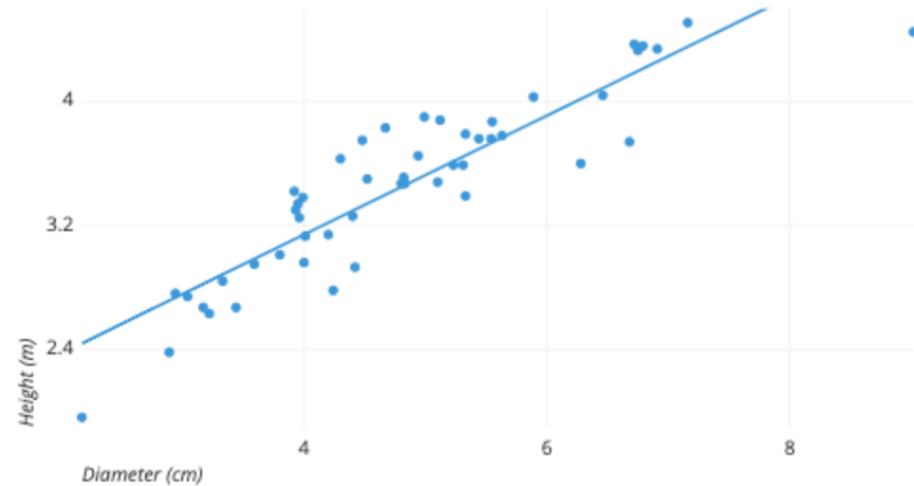
null / no relationship



strong, non-linear

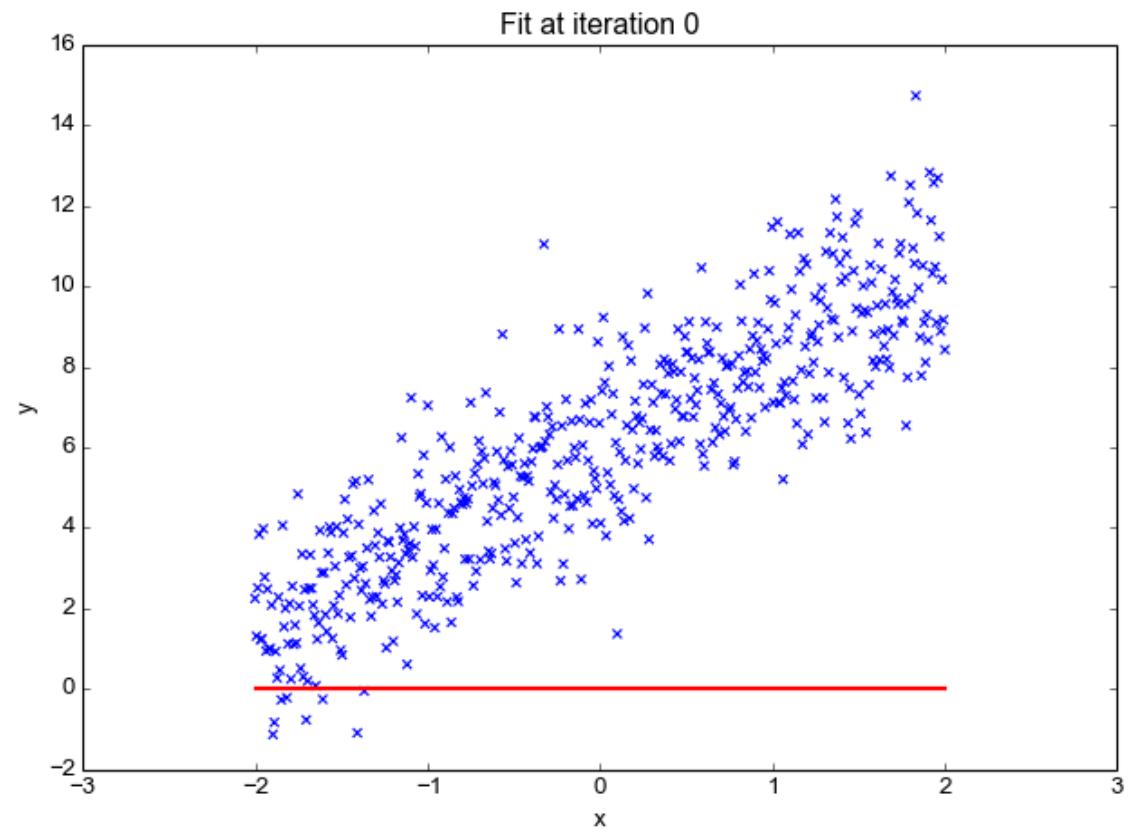
Let's add a trend line to the scatterplot

- When a scatter plot is used to look at a predictive or correlational relationship between variables, it is common to add a trend line to the plot showing the mathematically best fit to the data.
- This can provide an additional signal as to how strong the relationship between the two variables is, and if there are any unusual points that are affecting the computation of the trend line.



Let's visualize Linear Regression

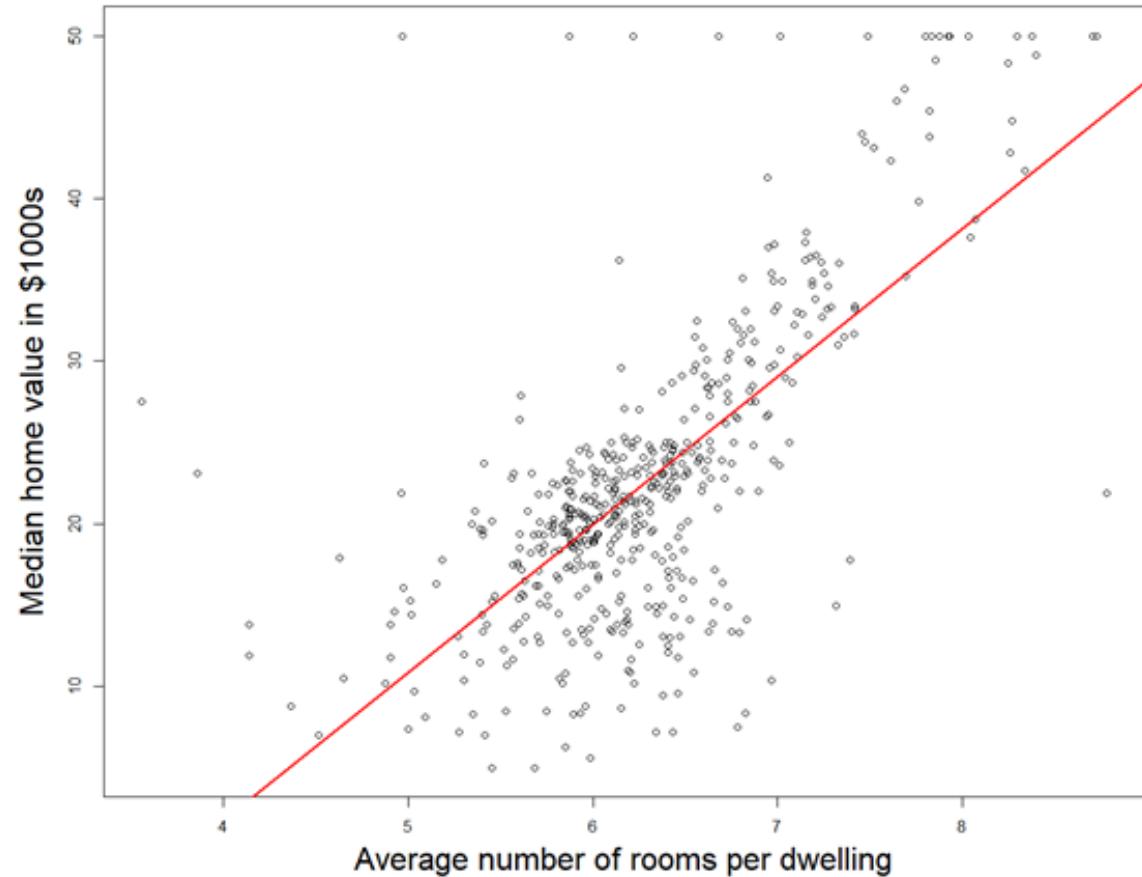
- Linear regression is a common statistical method, which has been adopted in machine learning and enhanced with many new methods for fitting the line and measuring error.



[Linear Regression - PRIMO.ai](#)

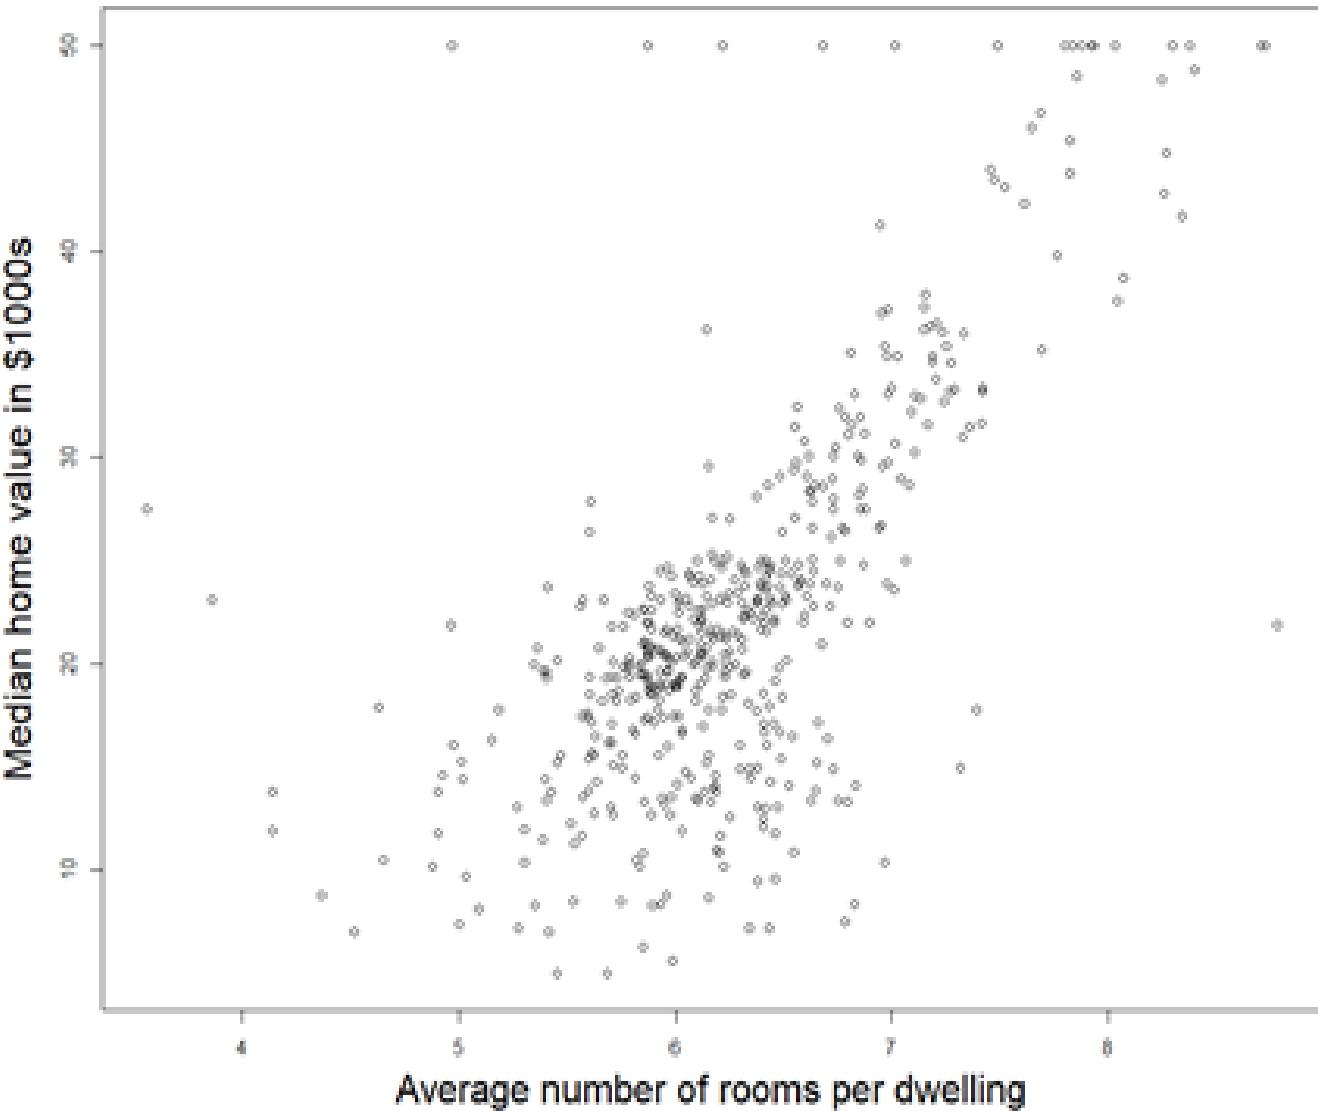
Formal definition of Regression

- **Regression analysis** is a set of statistical processes for estimating the relationships among variables.
- It includes many techniques for modeling and analyzing several variables, when the focus is on the relationship between an outcome and one or more predictors.



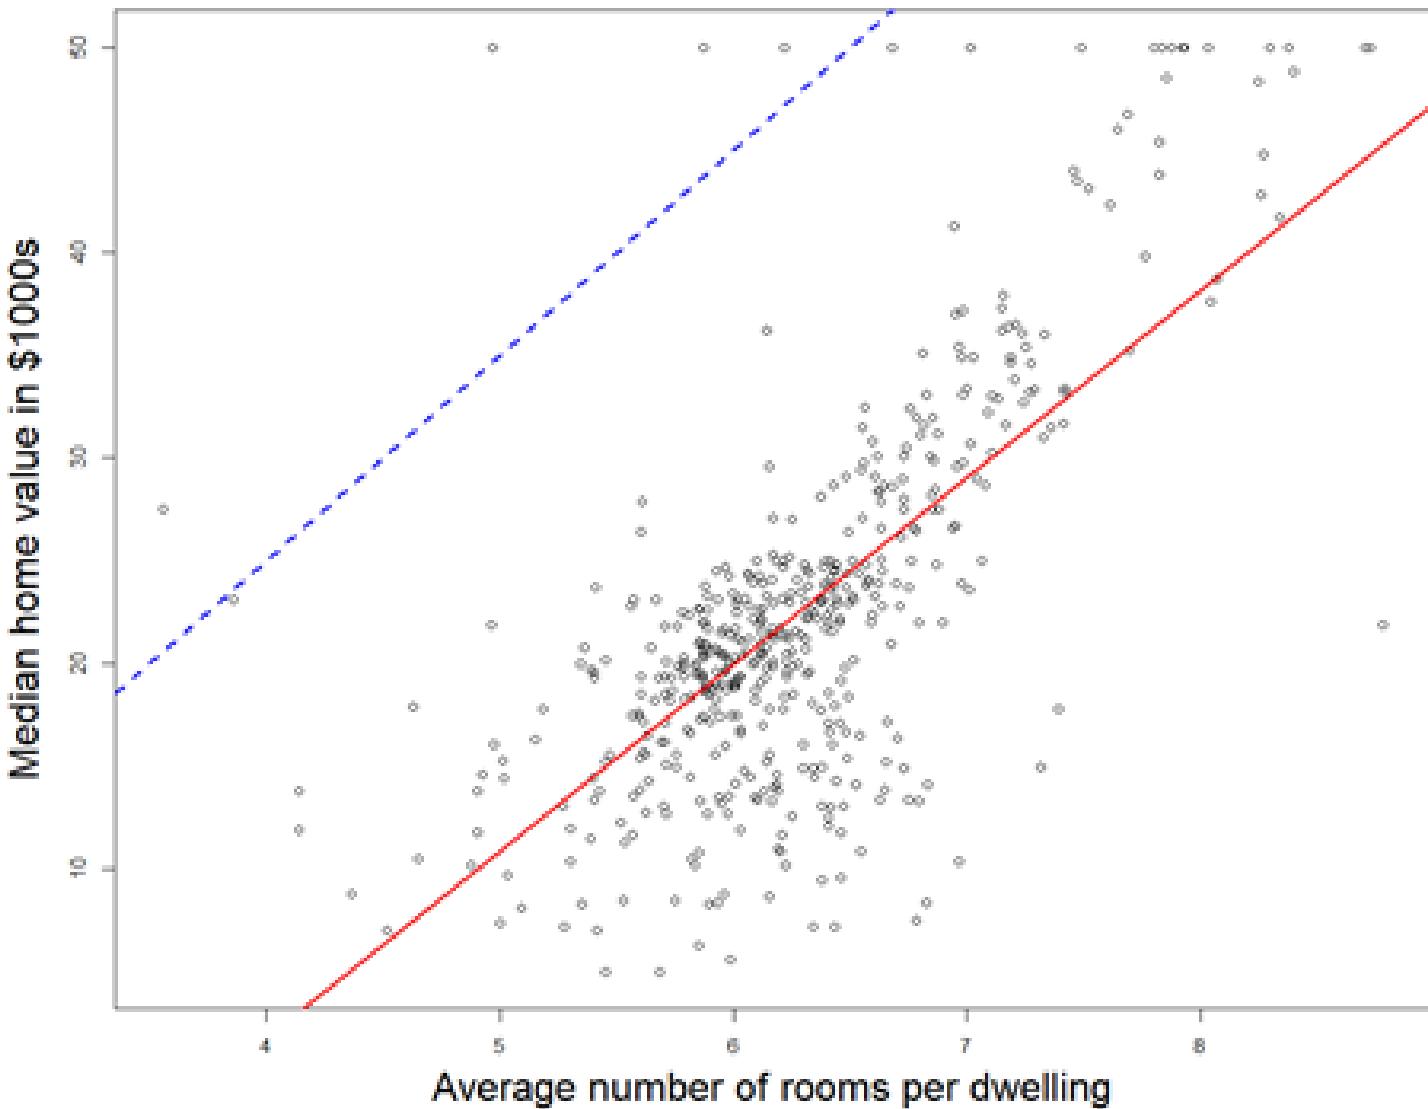
Simple Linear Regression

- Predict a quantitative outcome Y on the basis of a single predictor variable X .
- Maybe you are familiar with this: " $Y = mx + b$ ".
- Mathematically:
 $E[Y|X] = \beta_0 + \beta_1 X$ or $Y = \beta_0 + \beta_1 X + \varepsilon$
- Median home value = $\beta_0 + \beta_1 * \text{Average number of rooms} + \varepsilon$.



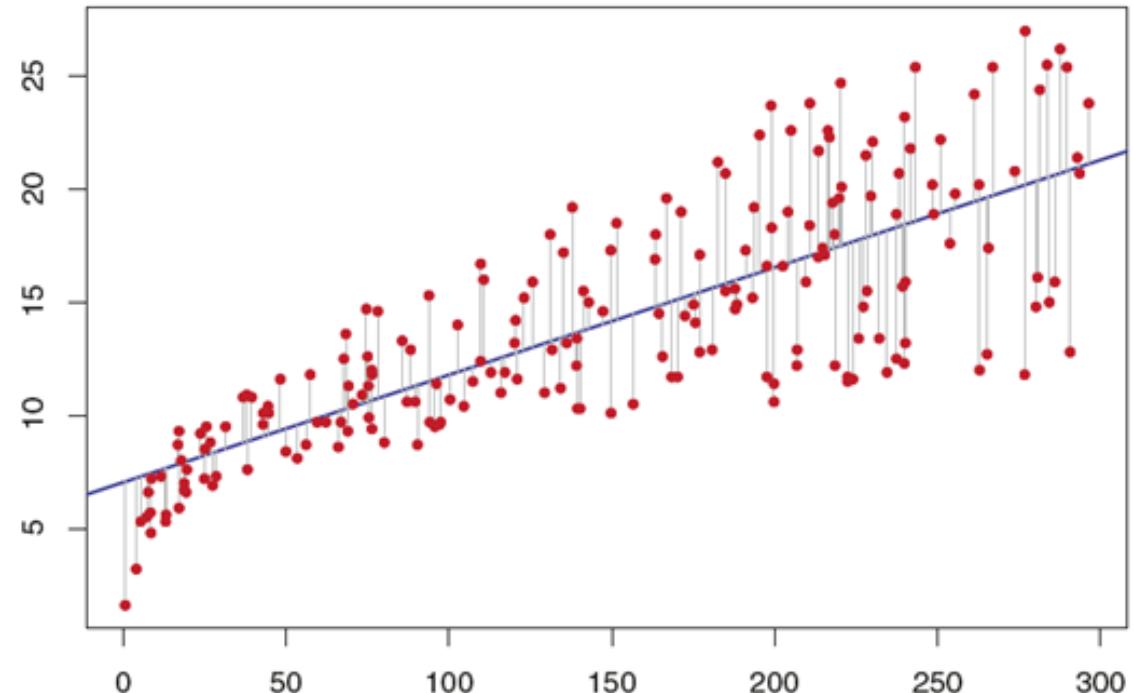
Estimate the coefficients

- In our model $Y = \beta_0 + \beta_1 X + \epsilon$
Here
 β_0 : "intercept", the expected value of Y when $X = 0$.
 β_1 : "slope", the average increase in Y associated with a one-unit increase in X .
 ϵ : "error term" or "residual".
- Which line fits the data better, blue or red?
We need to calculate β_0 and β_1 , so that the resulting line is as close as possible to these data points.



Estimate the coefficients

- There are a number of ways to define closeness. The most common approach involves minimizing the “residual sum of squares (RSS)”.
 - Intuition: minimize the total squared vertical distance of each data point to a potential linear line.
 - $\text{RSS} = \text{vertical distance}_1^2 + \text{vertical distance}_2^2 + \dots$
 - $\text{RSS} = (y_1 - \beta_0 - \beta_1 X_1)^2 + (y_2 - \beta_0 - \beta_1 X_2)^2 + \dots$



R or Python can be used to do this!

- Our linear model: $E[\text{Median home value}] = -34.7 + 9.1 \times \text{Average number of rooms.}$
- Interpretation: associated with a one-unit increase in average number of rooms, the median home value is estimated to increase by \$9,100.

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-34.671	2.650	-13.08	<2e-16 ***
rm	9.102	0.419	21.72	<2e-16 ***

Goodness of fit: R-square

- R-squared statistic: measures the proportion of variability in Y that can be explained using X.
- It ranges from 0 to 1 (can be negative in some rare cases).
- Higher is better, but it can be challenging to determine what is good.
- In our case, R-squared = 0.48.

```
> fit1 <- lm(medv ~ rm, data = Boston)
> summary(fit1)

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -34.671     2.650  -13.08  <2e-16 ***
rm            9.102     0.419   21.72  <2e-16 ***
---
Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
Residual standard error: 6.616 on 504 degrees of freedom
Multiple R-squared:  0.4835,    Adjusted R-squared:  0.4825
F-statistic: 471.8 on 1 and 504 DF,  p-value: < 2.2e-16
```

Assumptions when building a model

- Linearity: is the relationship between outcome and predictors really linear?
- Normality: does the error term follow a Normal distribution?
- Equal Variance: does the error term have equal variance?
- Independence of observations: does your observations intendent from each other? Eg. can we assume one observation is independent from another observation?

Try it out!

- [Linear Regression Example — scikit-learn 0.24.2 documentation](#)

The screenshot shows a Jupyter Notebook interface with the following details:

- File Bar:** File, Edit, View, Run, Kernel, Tabs, Settings, Help.
- Launcher:** plot_ols.ipynb
- Code Cell:** [1]: %matplotlib inline
- Title:** Linear Regression Example
- Description:** The example below uses only the first feature of the `diabetes` dataset, in order to illustrate the data points within the two-dimensional plot. The straight line can be seen in the plot, showing how linear regression attempts to draw a straight line that will best minimize the residual sum of squares between the observed responses in the dataset, and the responses predicted by the linear approximation.
- Note:** The coefficients, residual sum of squares and the coefficient of determination are also calculated.
- Code:**

```
%matplotlib inline

# Code source: Jaques Grobler
# License: BSD 3 clause

import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score

# Load the diabetes dataset
diabetes_X, diabetes_y = datasets.load_diabetes(return_X_y=True)

# Use only one feature
diabetes_X = diabetes_X[:, np.newaxis, 2]

# Split the data into training/testing sets
diabetes_X_train = diabetes_X[:-20]
diabetes_X_test = diabetes_X[-20:]

# Split the targets into training/testing sets
diabetes_y_train = diabetes_y[:-20]
diabetes_y_test = diabetes_y[-20:]
```

Classification



Data for Making Predictions

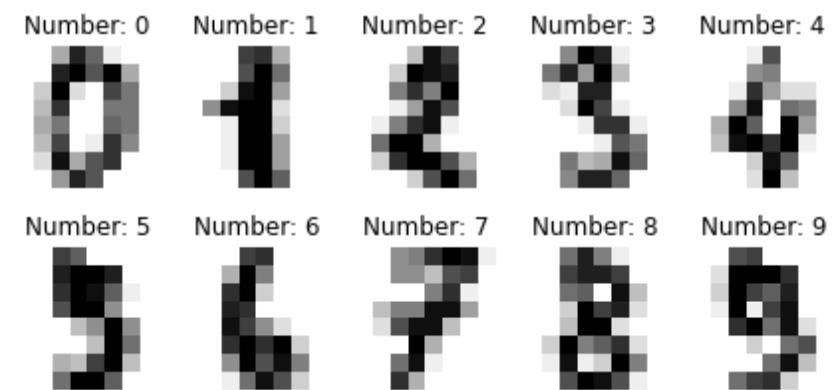
- When using statistics or machine learning to make predictions, we allow the model to learn based on labeled examples.
- A **label** is what we're predicting
- A **feature** is an input variable
- A **labeled example** includes feature(s) and the label.
- A dataset consists of many labeled examples.

Preparation of Data

- Try to get all numeric features on the same scale (usually, scale to between 0 and 1)
 - For example, compare the scale of house prices with the scale of square footage? (\$100,000 vs 1200 sqft)
- Map categorical features to integers
 - For example, {apartment, condo, house} -> {0, 1, 2}
- Reserve some data to evaluate the model – called a train/test split
 - Model sees train data during training phase
 - Model is evaluated on test data after training – to show performance on unseen data

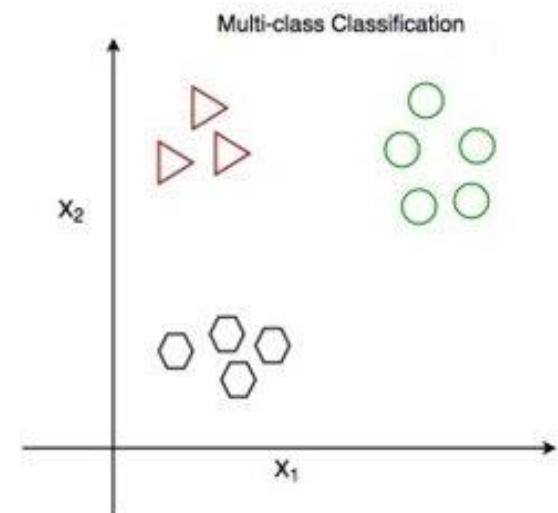
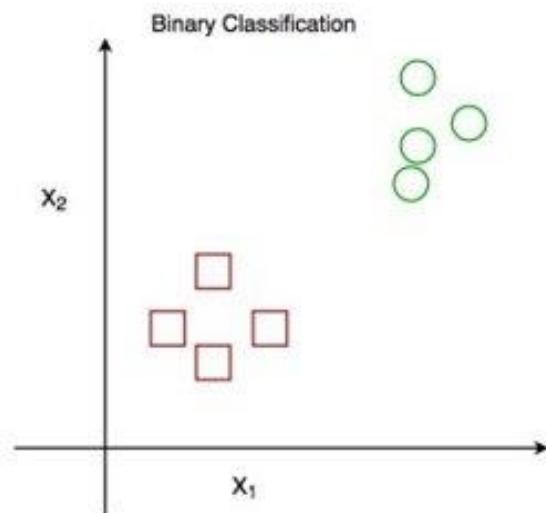
Defining the Classification Problem

- Learn how to make predictions from past examples
- What are the predictions? Finite set of class labels or categories
 - Given medical history, does an individual have an illness or not? (2 classes)
 - Given past weather information, will the weather tomorrow be sunny, rainy, or cloudy? (3 classes)
 - For each image, what digit 0-9 is represented? (10 classes)



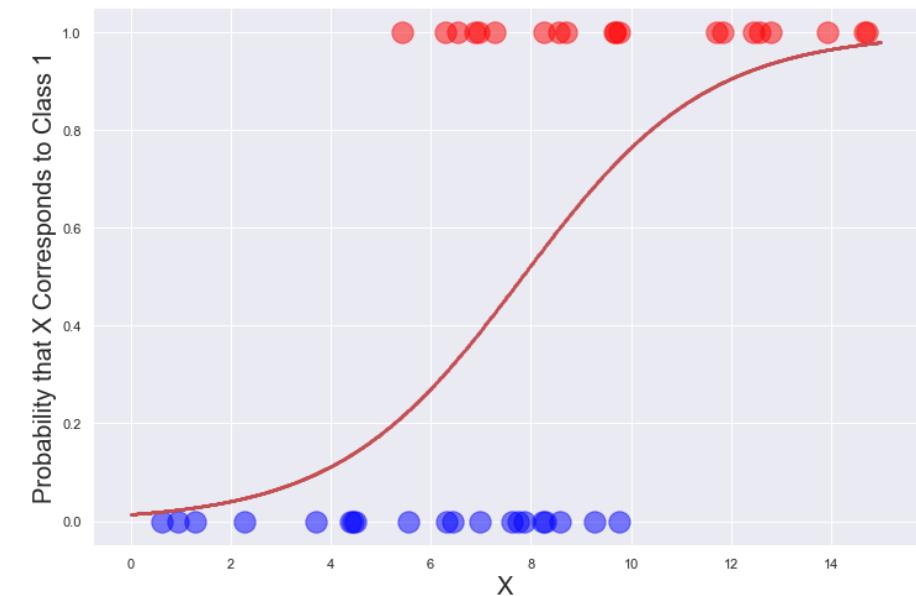
Binary Classification

- Classification where two classes are being predicted:
 - Is this email SPAM or NOT SPAM?
 - Is this sentence POSITIVE or NEGATIVE?
 - Given medical history, does an individual HAVE AN ILLNESS or NOT?



Binary Classification Approach

- Map one class to 0 (NOT SPAM) and the other class to 1 (SPAM)
- Model produces a score between 0 and 1
 - If this score is ≤ 0.5 , classify example as NOT SPAM
 - If this score is > 0.5 , classify example as SPAM
- Different methods for model to produce this score are different machine learning algorithms and approaches



Evaluating a Binary Classification Model

How do we know if our model is performing well?

$$\text{Accuracy} = \frac{\text{Number of Correct predictions}}{\text{Total number of predictions made}}$$

		ACTUAL	
		Positive	Negative
PREDICTED	Positive	TRUE POSITIVE	FALSE POSITIVE
	Negative	FALSE NEGATIVE	TRUE NEGATIVE

*Confusion Matrix –
allows you to visualize performance*

Try it out!



[Google CoLab Notebook](#)

AI in Industry

Overview

Topics for the day

- Introduction to AI in Industry
- Natural Language Processing
- Computer Vision
- Q&A

Speakers

- Daisy Deng – Natural Language Processing
- Max Kaznady – Computer Vision

AI by Industry

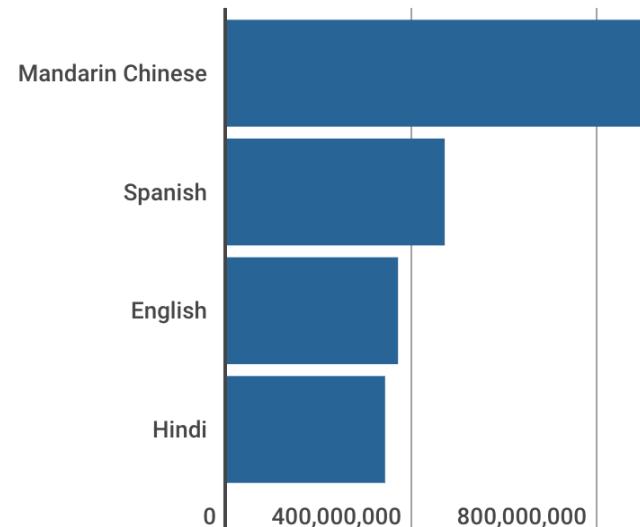


Natural Language Processing



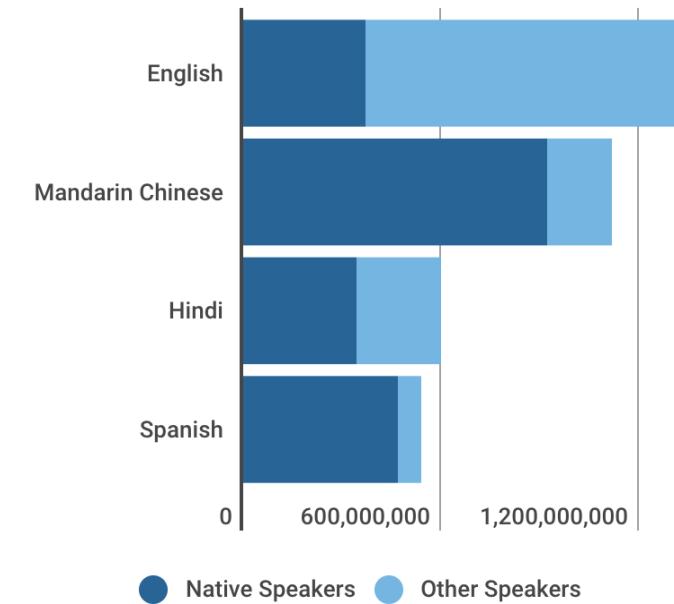
What's the
most spoken
language?

Languages with the most native speakers



Mandarin Chinese is the largest language in the

Languages with the most speakers



How many
number of
languages are
there in the world?



The number of natural languages in the world



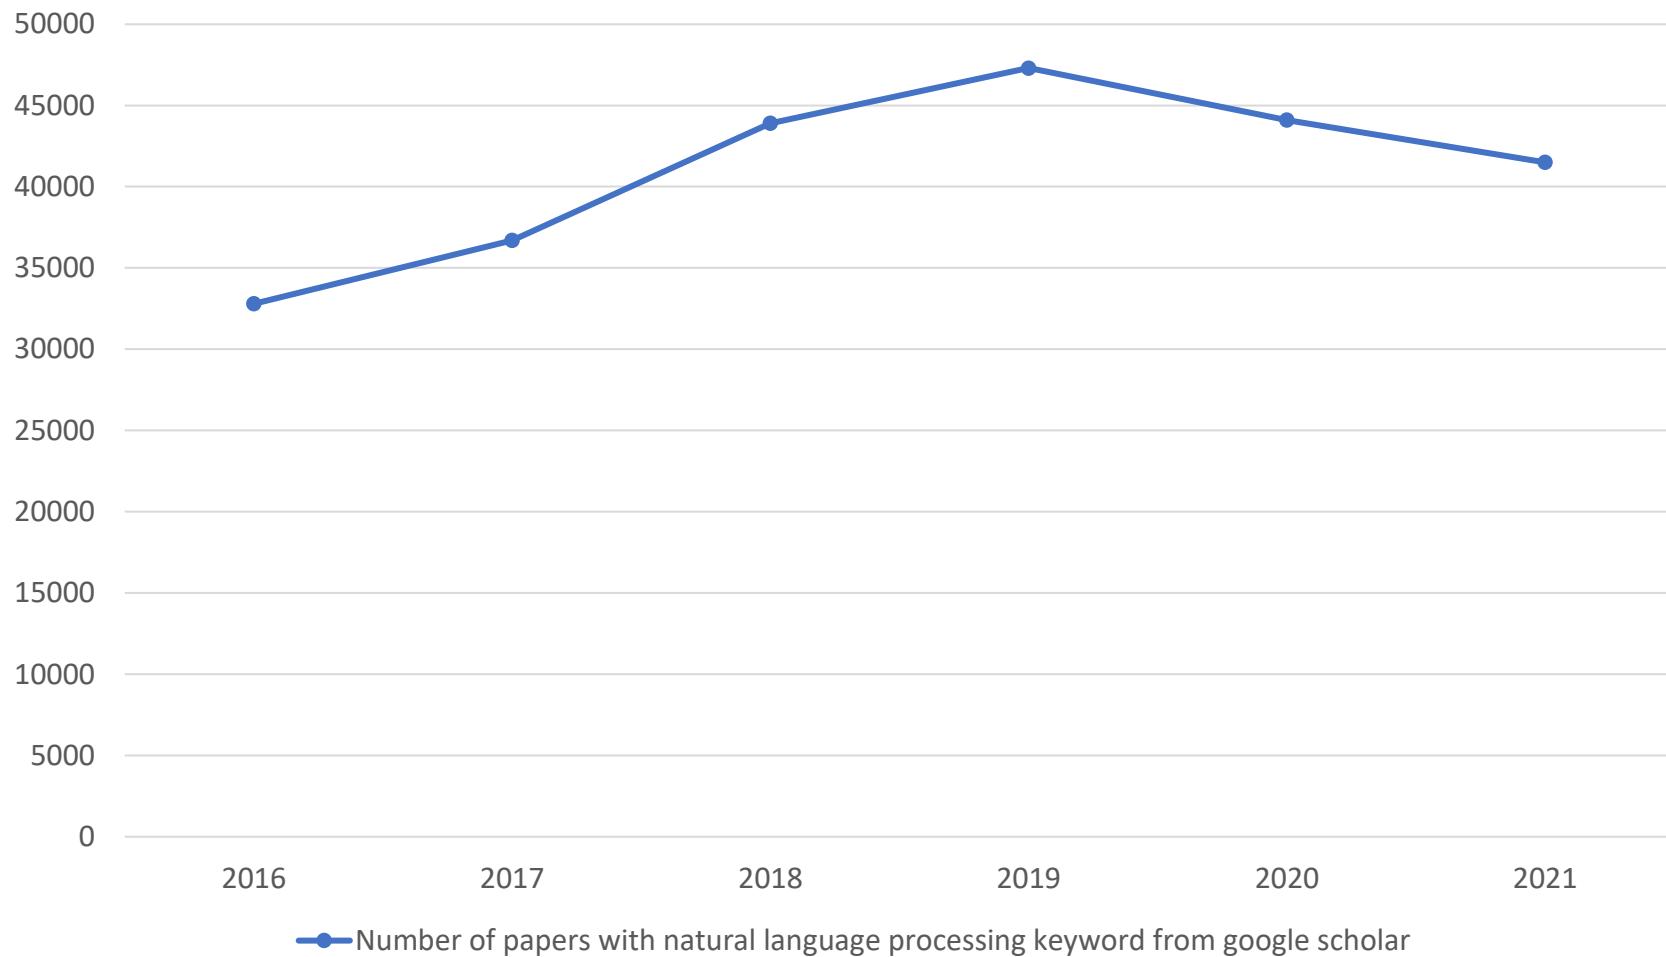
4,065 have a developed writing system.

7,139 languages are spoken today

<https://www.ethnologue.com/guides/how-many-languages>

<https://www.ethnologue.com/browse/countries>

Number of papers with natural language processing keyword from google scholar



What is Natural Language Processing NLP

Wiki definition

- subfield of [linguistics](#), [computer science](#), and [artificial intelligence](#) concerned with the interactions between computers and human language.
- Extract meaningful information from natural language input
- Produce useful natural language output based in the input

Modern Language applications

- Email categorization (spam, promotion, primary etc)
- Search engine (Bing, Google ...)
- Chat bot (MSFT Cortana, Google Assistant, Apple Siri)
- Translation systems (Bing Microsoft Translator, Google Translate)
- Brand sentiment monitoring
- Customer review analysis
- News Digest
- Summarize for a 2nd grader (<https://beta.openai.com/examples/default-summarize>)
- ...

NLP Applications

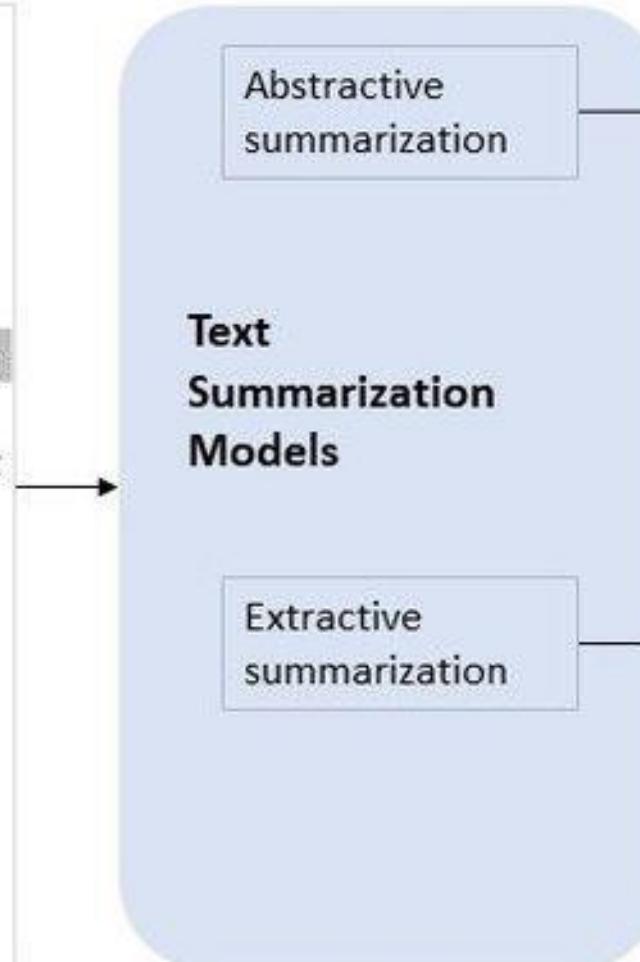
- Classification:
 - Best burger ever.
 - Service is disappointing.
- Named entity recognition
 - Mary have a doctor's appointment at 11AM EDT in xxx Massachusetts Ave, Cambridge.

Entities: Person, event, time, location
- Translation:
 - Input: 我的名字叫黛西. (Simplified Chinese)
 - Output: My name is Daisy (English)

NLP Applications (Summarization)

Input Article

Marseille, France (CNN) The French prosecutor leading an investigation into the crash of Germanwings Flight 9525 insisted Wednesday that he was not aware of any video footage from on board the plane. Marseille prosecutor Brice Robin told CNN that " so far no videos were used in the crash investigation . " He added, " A person who has such a video needs to immediately give it to the investigators . " Robin\ 's comments follow claims by two magazines, German daily Bild and French Paris Match, of a cell phone video showing the harrowing final seconds from on board Germanwings Flight 9525 as it crashed into the French Alps . All 150 on board were killed. Paris Match and Bild reported that the video was recovered from a phone at the wreckage site. ...



Generated summary

Prosecutor : " So far no videos were used in the crash investigation "

Extractive summary

marseille prosecutor brice robin told cnn that " so far no videos were used in the crash investigation . " robin \ 's comments follow claims by two magazines , german daily bild and french paris match , of a cell phone video showing the harrowing final seconds from on board germanwings flight 9525 as it crashed into the french alps . paris match and bild reported that the video was recovered from a phone at the wreckage site .

NLP Applications (Continued)

- Question and Answering (extractive and abstractive)
 - Context:
 - Question
 - Answer
- More....

Passage Sentence

In meteorology, precipitation is any product of the condensation of atmospheric water vapor that falls under gravity.

Question

What causes precipitation to fall?

Answer Candidate

gravity

NLP Applications Pop Quiz

On a piece of tech/finance news, you want to find all the companies the news mentioned, what application would you use to solve the problem?

1. Classification
2. Question and Answering
3. Summarization
4. Named Entity Recognition

NLP Linguistic Tasks

- Tokenization: split sentences to words
 - “this is a simple sentence” -> [“this”, “is”, “a”, “simple”, “sentence”]
- Stop words removal: remove high-frequency words with little or no semantic value
 - “This is a simple sentence” -> “This simple sentence”
- Stemming: get the root form of a word
 - Fish, fishing, fishes -> fish
 - Better -> better
- Lemmatization: remove inflectional endings, return the base dictionary form of a word
 - Gone, going, went -> go
 - Good, better, best-> good

NLP Linguistic Tasks

- Part-of-speech Tagging (POS)
 - This is a simple sentence -> DT VBS DT JJ NN
- Dependency parsing
 - Syntax structure
- Constituency Parsing
 - Phrase structure

Movie Review Sentiment Analysis

Task: Given a piece of movie review text, determine if the review is positive or negative.

What NLP application would you use?

Dataset: Large Movie Review Dataset

- **<https://ai.stanford.edu/~amaas/data/sentiment/>**
- highly polar movie reviews
- 25,000 for training
- 25,000 for testing.

Data Cleaning/Text preprocessing

- Remove irrelevant parts.
- Tokenize the input
- Remove Stop Words
- Lemmatization

Convert text to features in classification model

- Question: you have learned classification models in the previous session. How would you translate the text into the numerical inputs for a classification model?

Split text to words and counting words!

Text vectorization: given a dictionary, one-hot encoding

N-gram model

- Multiple word is considered as a unit for data processing
- S1 = "Natural language processing is interesting."
S2 = "Learning a programming language is interesting"

Uni-gram: Vocabulary = set("a", "interesting", "is", "natural", "language", "learning", "processing", "programming")

- Size of vocabulary is 8
- S1 = [0, 1, 1, 1, 1, 0, 1, 0]
- S2 = [1, 1, 1, 0, 1, 1, 0, 1]

N-gram model (Continued)

- "I love vanilla, but I hate chocolate"
- "I love chocolate, but I hate vanilla"

With Uni-gram model, order info is completely discarded

```
vectorizer = CountVectorizer(ngram_range=(1,1))
sentence = "I love chocolate, but I hate vanilla"
X = vectorizer.fit_transform([sentence])
print(f"features for \"{sentence}\" is {vectorizer.get_feature_names()} ")
sentence = "I love vanilla, but I hate chocolate"
X = vectorizer.fit_transform([sentence])
print(f"features for \"{sentence}\" is {vectorizer.get_feature_names()} ")

features for "I love chocolate, but I hate vanilla" is ['but', 'chocolate', 'hate', 'love', 'vanilla']
features for "I love vanilla, but I hate chocolate" is ['but', 'chocolate', 'hate', 'love', 'vanilla']
```

N-gram model (continued)

- S1 = "Natural language processing is interesting."
S2 = "Learning a programming language is interesting"

Bi-gram: Vocabulary = set("natural language", "language processing", "processing is", "is interesting", "learning a", "a programming", "programming language", "language is")

- Size of vocabulary is 8
- S1 = [1, 1, 1, 1, 0, 0, 0, 0]
- S2 = [0, 0, 0, 1, 1, 1, 1, 1]

N-gram model (Continued)

```
vectorizer = CountVectorizer(ngram_range=(2,2))
sentence = "I love chocolate, but I hate vanilla"
X = vectorizer.fit_transform([sentence])
print(f"features for \"{sentence}\" is {vectorizer.get_feature_names()} ")
sentence = "I love vanilla, but I hate chocolate"
X = vectorizer.fit_transform([sentence])
print(f"features for \"{sentence}\" is {vectorizer.get_feature_names()} ")

features for "I love chocolate, but I hate vanilla" is ['but hate', 'chocolate but', 'hate vanilla', 'love chocolate']
features for "I love vanilla, but I hate chocolate" is ['but hate', 'hate chocolate', 'love vanilla', 'vanilla but']
```

With Bi-gram model, some sequence info is
kept

Document Frequency

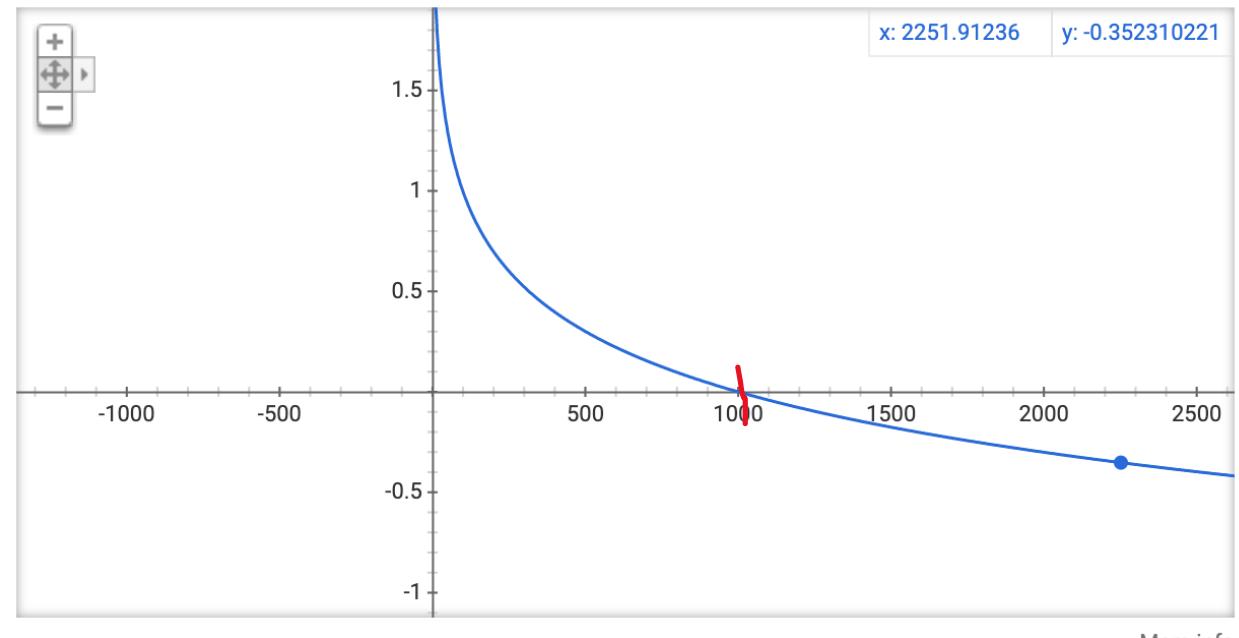
- Common words for movie review: director, movie, story, film
- Words that differentiate: worst, best, wonderful, awful
- The **document frequency** of a word is equal to the number of documents which contain the word divided by the total number of documents.
 - Total 100 reviews
 - document frequency of word “movie” might be close to 1
 - Document frequency of word “awful” probably much less than 1

Inverse Document Frequency

$$idf(w, D) = \log\left(\frac{1 + |D|}{1 + df(w, D)}\right)$$

- Question: for words that are not occurred in the document, it would have a very high IDF. How can we mitigate this?

Graph for $\log((1+1000)/(1+x))$



TF-IDF



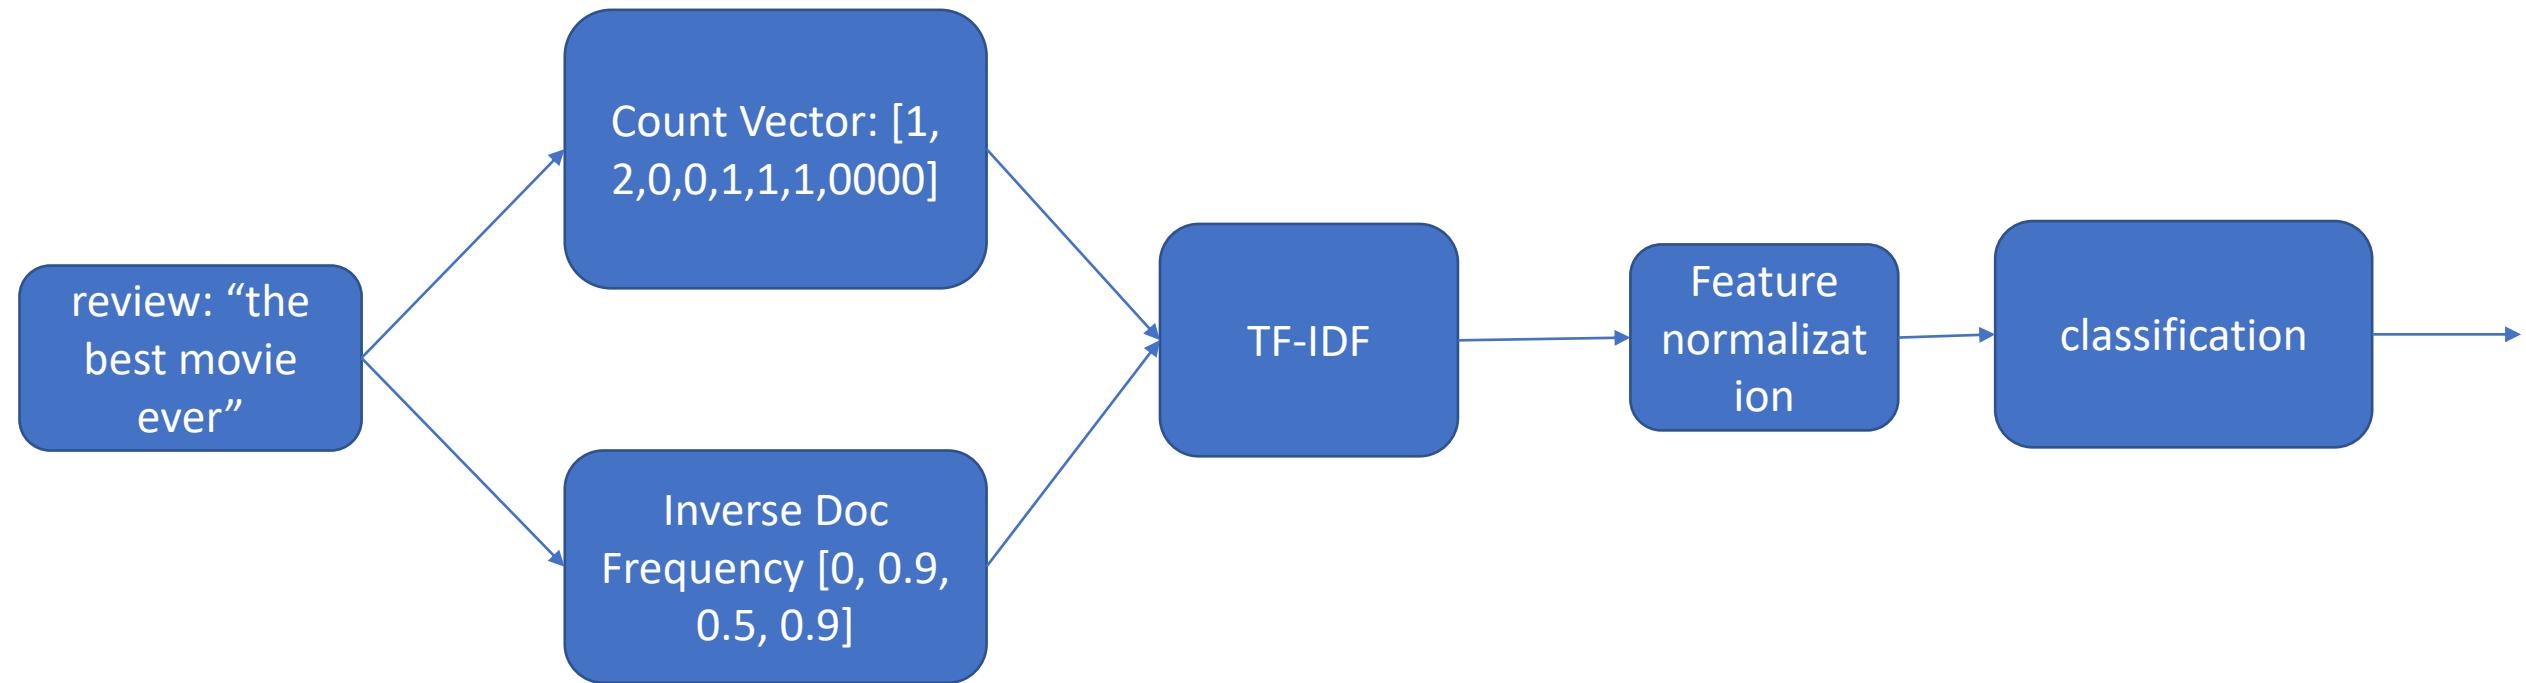
Term Frequency – Inverse Document Frequency

$tf(w, d) = f_d(w)$: frequency of w in document d

$$idf(w, D) = \log\left(\frac{1 + |D|}{1 + df(w, D)}\right)$$

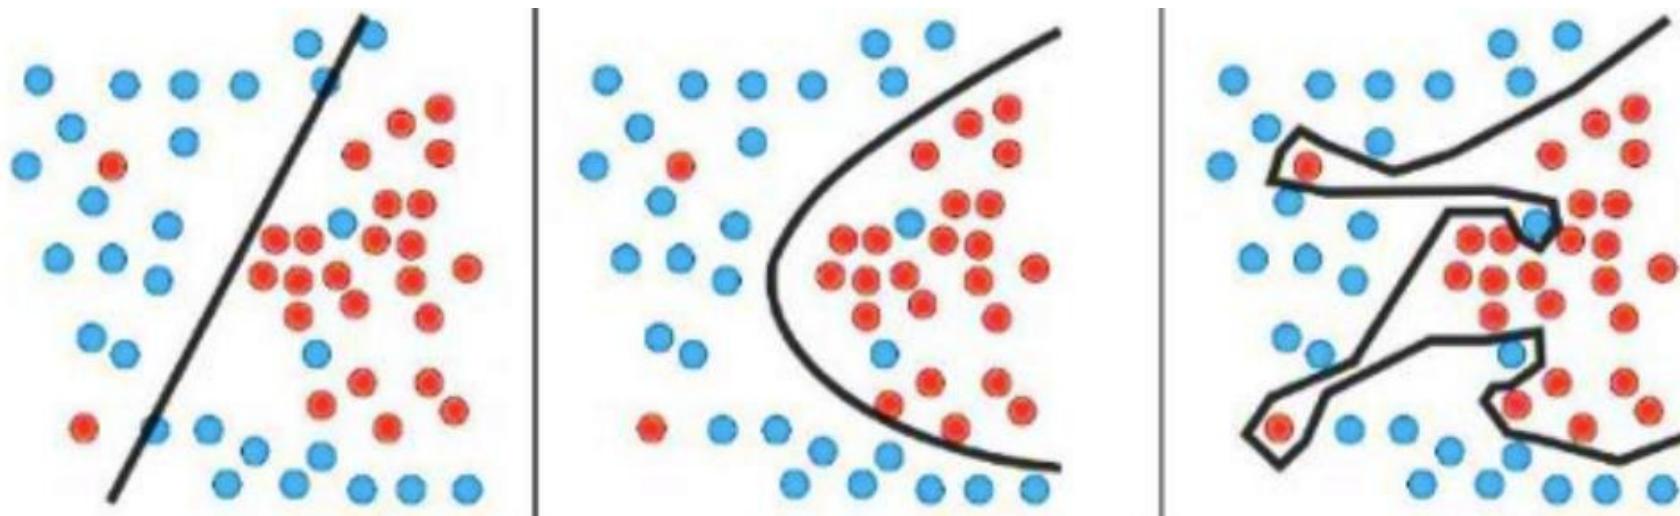
$$tfidf(w, d, D) = tf(w, d) \times idf(w, D)$$

Putting it together



What N should be in N-gram model?

Model Complexity vs Model Quality



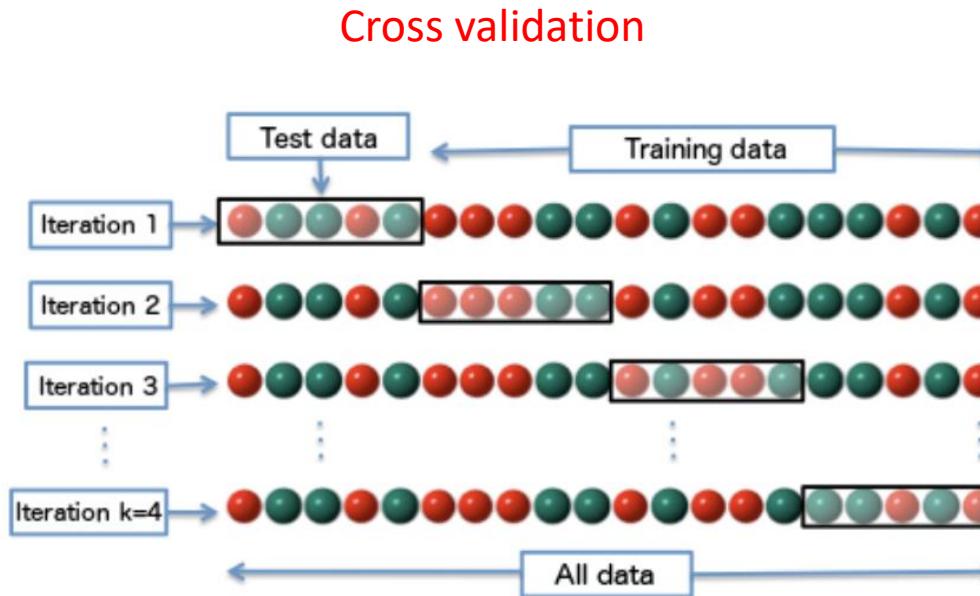
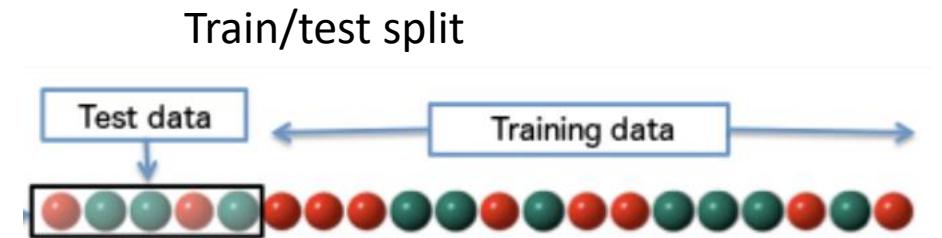
<https://towardsdatascience.com/techniques-for-handling-underfitting-and-overfitting-in-machine-learning-348daa2380b9>

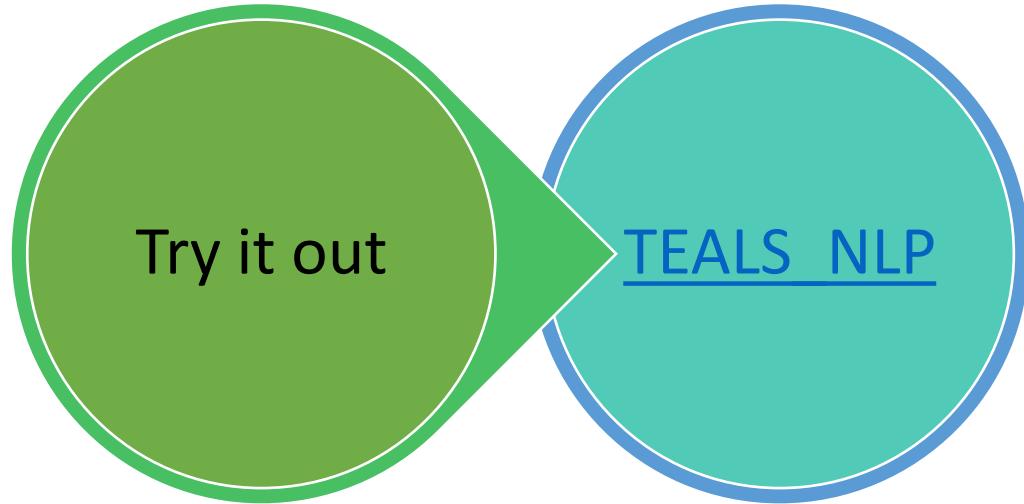
<https://towardsdatascience.com/understanding-the-bias-variance-tradeoff-165e6942b229>

Measure Model Quality/Performance

Is Bi-Gram better than
Uni-Gram?

Can the Bi-gram model
perform consistently well
on unseen data?

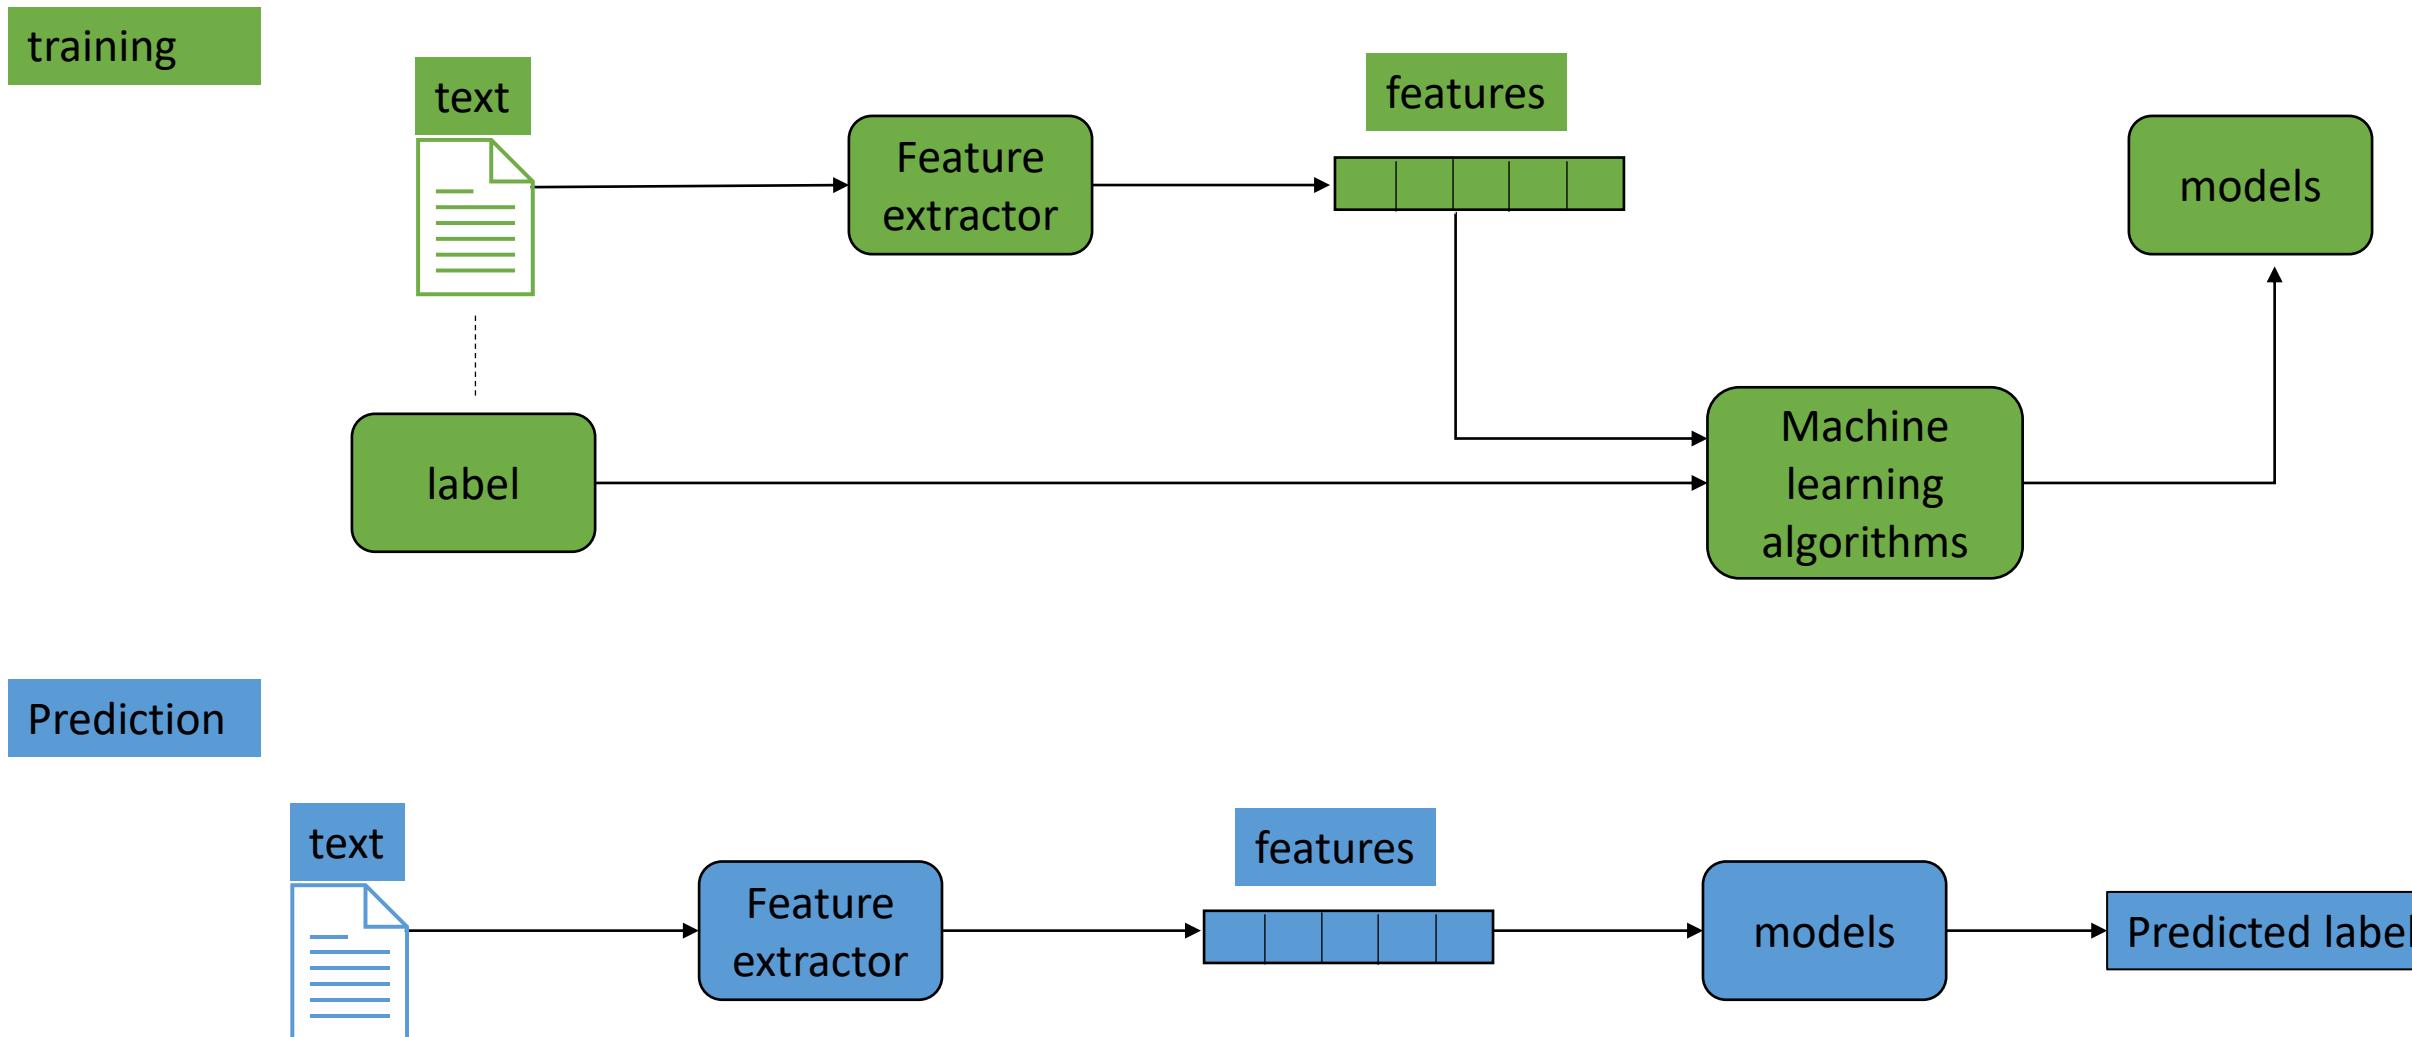




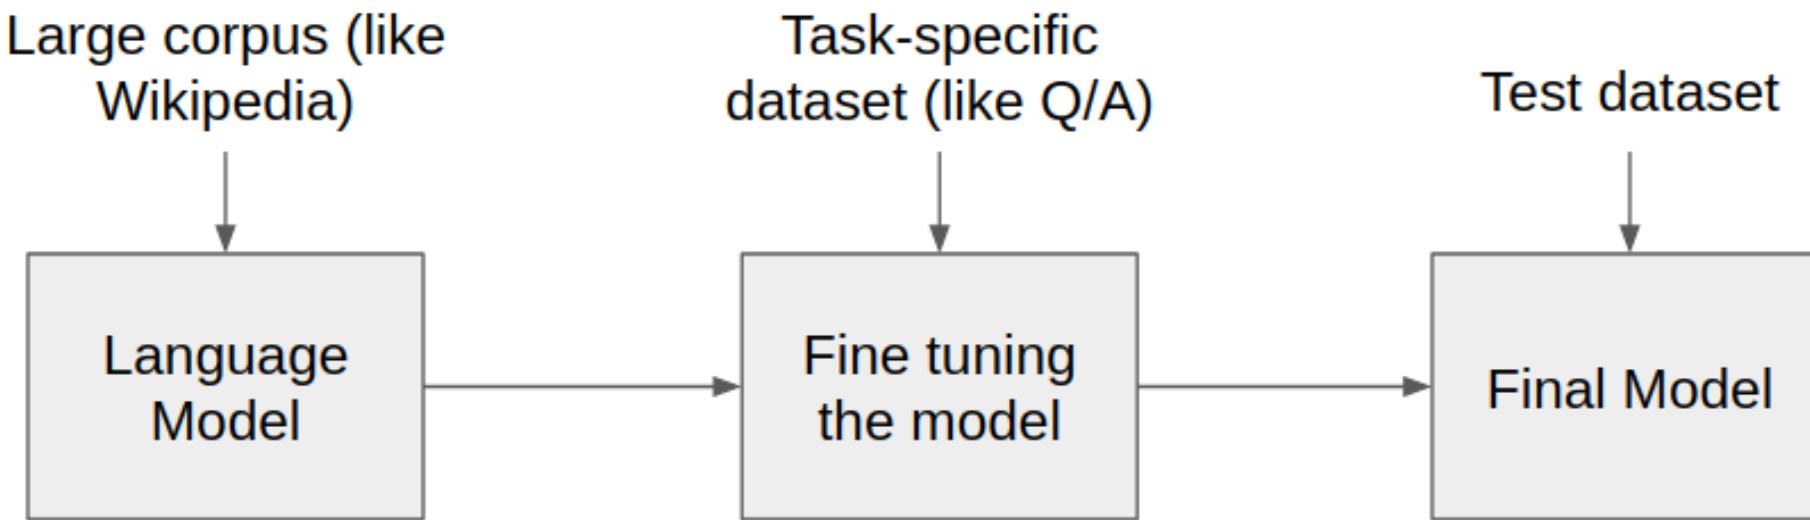
Try it out

TEALS_NLP

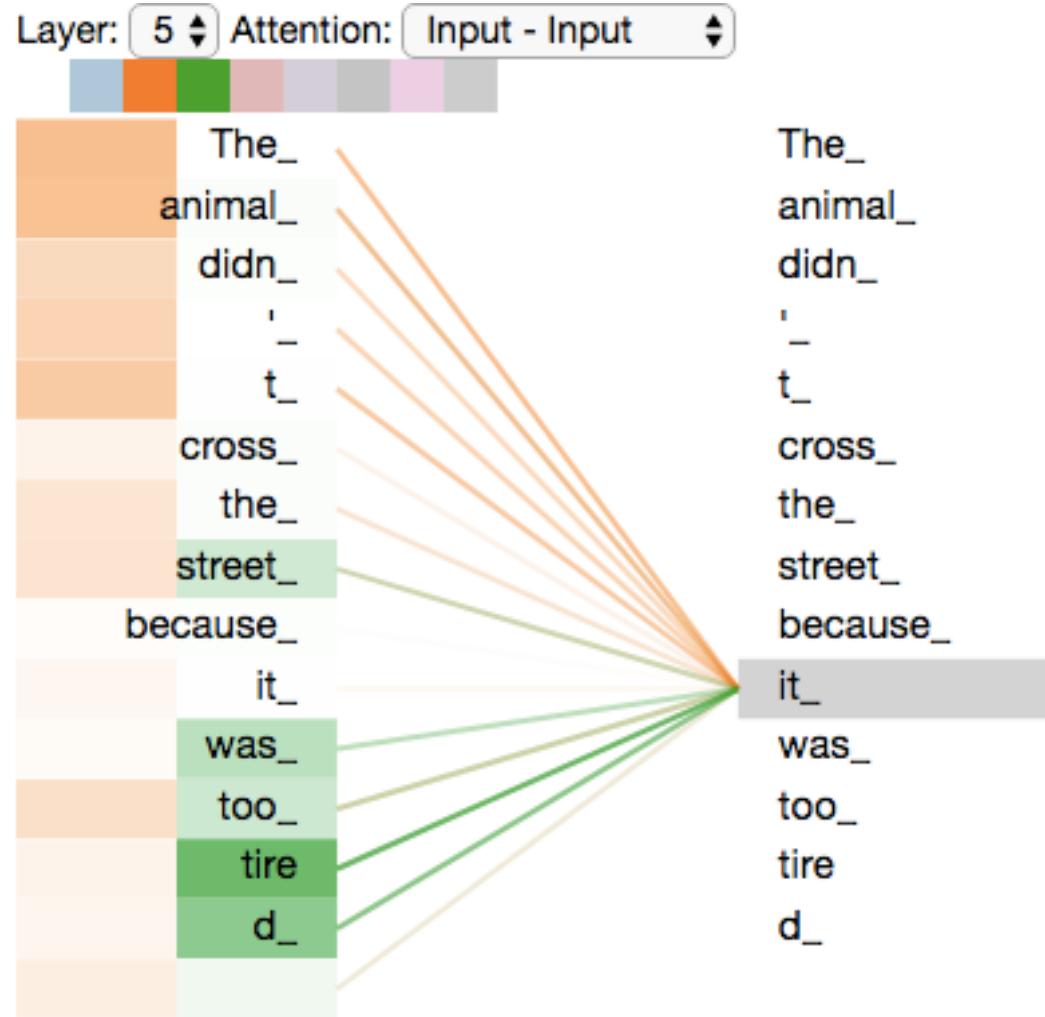
NLP for Classification



Transfer Learning



BERT (Bidirectional Encoder Representations from Transformers)



BERT (Cont'd)

- Encode information with longer dependencies from both directions
 - **340 million parameters**
 - More than 100 natural languages support
 - Much better performance with much less labeled data
-
- <https://www.kaggle.com/atulanandjha/bert-testing-on-imdb-dataset-extensive-tutorial>

GPT-3

- **Generative Pre-trained Transformer 3**
- **175 billion parameters**

Q: What is your favorite animal?

A: My favorite animal is a dog.

Q: Why?

A: Because dogs are loyal and friendly.

Q: What are two reasons that a dog might be in a bad mood?

A: Two reasons that a dog might be in a bad mood are if it is hungry or if it is hot.

Q: How many eyes does a giraffe have?

A: A giraffe has two eyes.

Q: How many legs does a frog have?

A: A frog has four legs.

Q: Are there any animals with three legs?

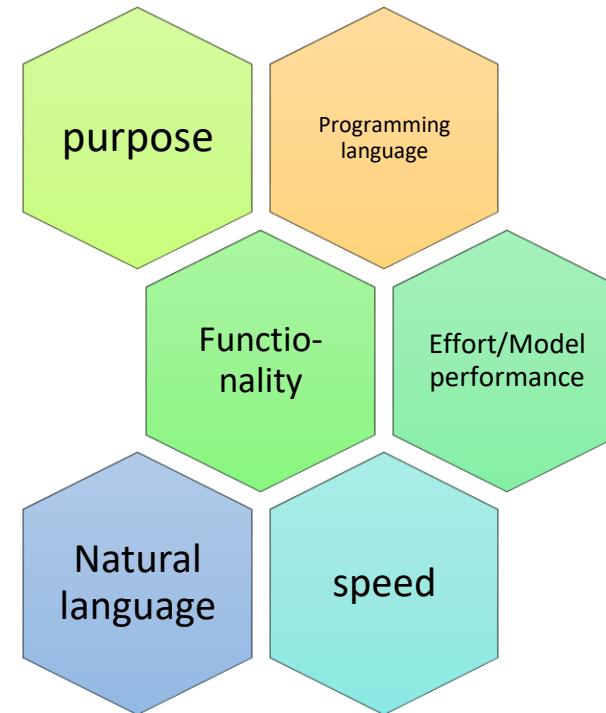
A: No, there are no animals with three legs.

Q: Why don't animals have three legs?

A: Animals don't have three legs because they would fall over.

Useful python packages for NLP

- [NLTK](#)
- [Spacy](#)
- [CoreNLP](#)
- [Huggingface](#)



Computer Vision



Computer Vision Tasks

Classification



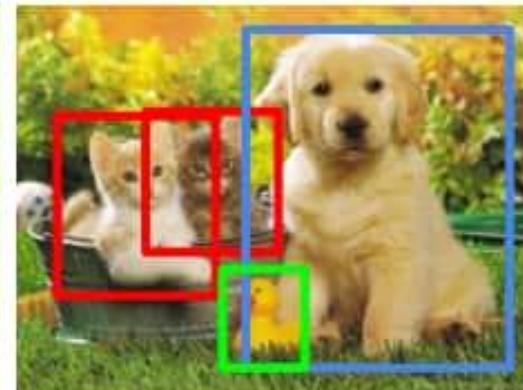
CAT

Classification + Localization



CAT

Object Detection



CAT, DOG, DUCK

Instance Segmentation



CAT, DOG, DUCK

Single object

Multiple objects

Image Segmentation flavors

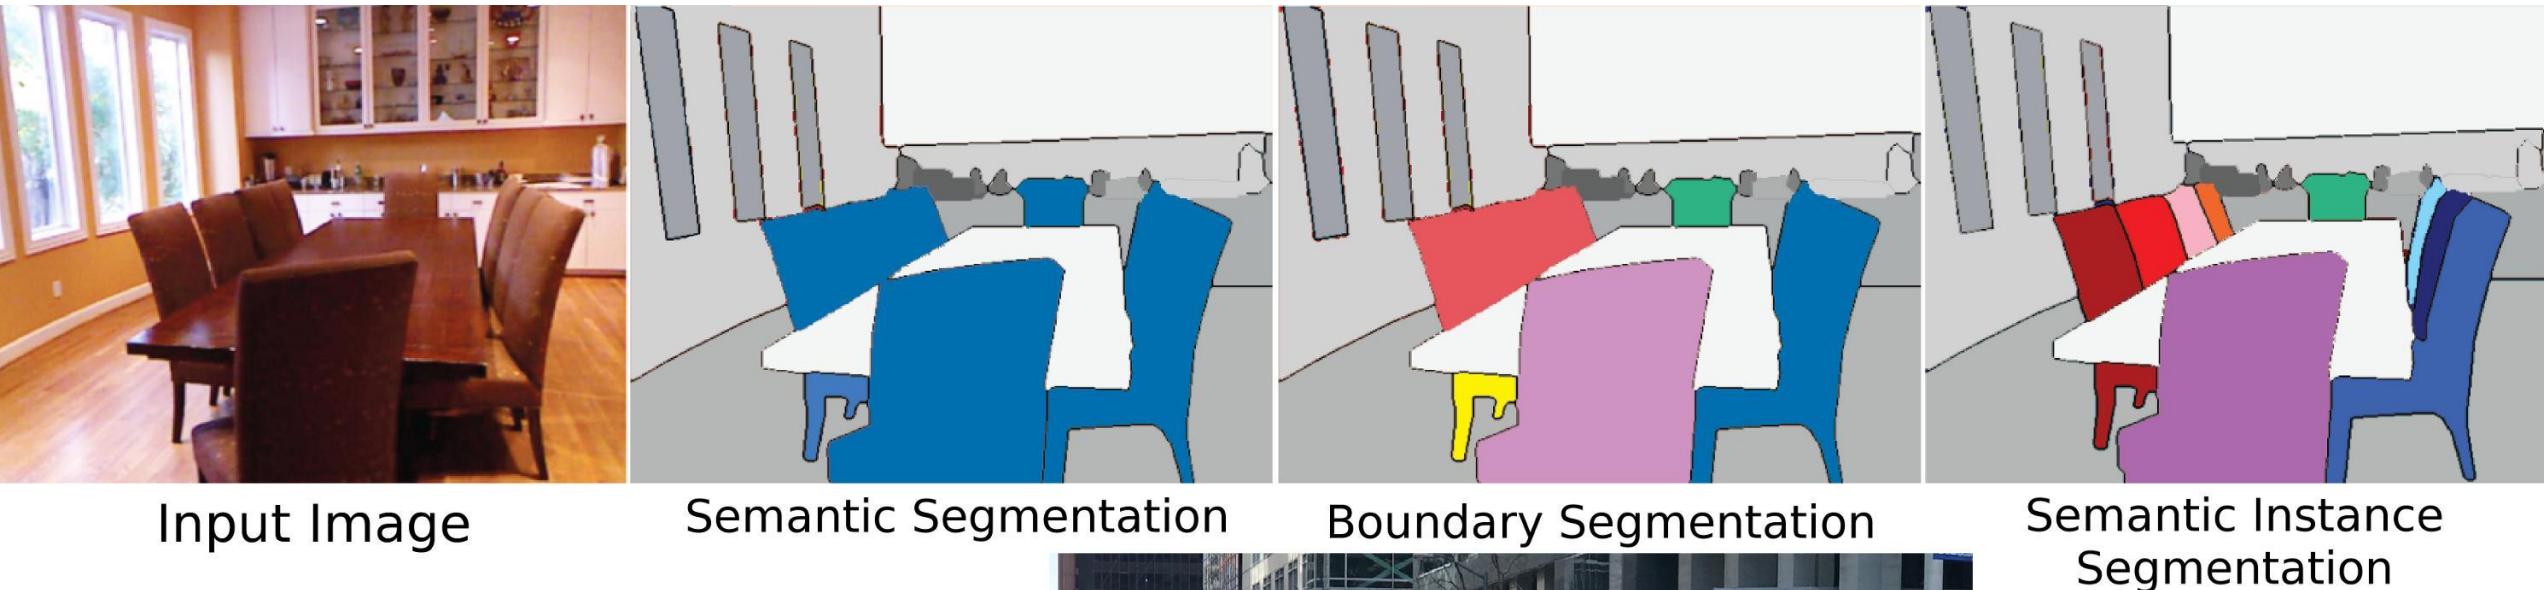
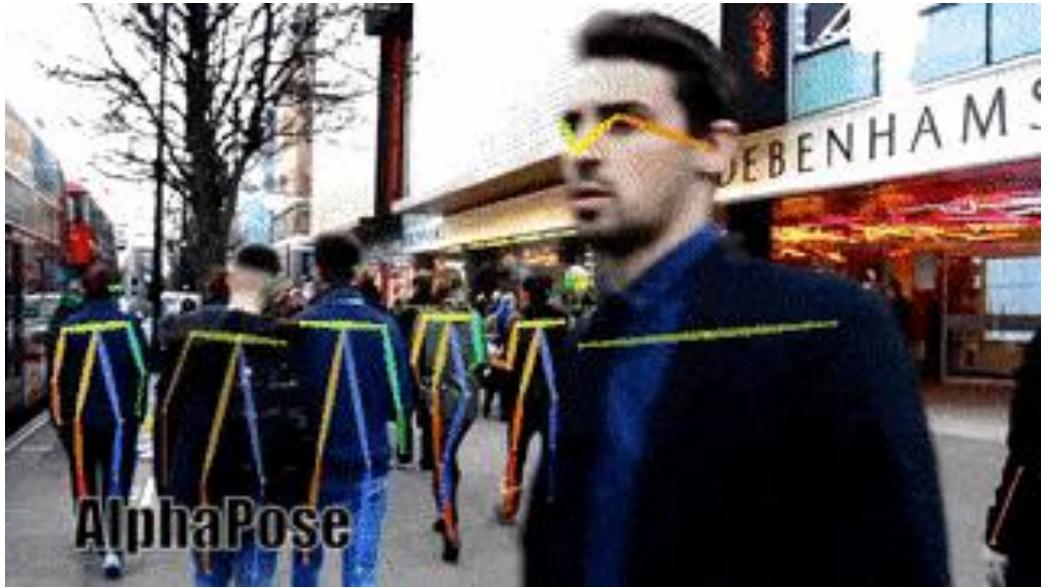


Image Masking

Boundary matters a lot in segmentation

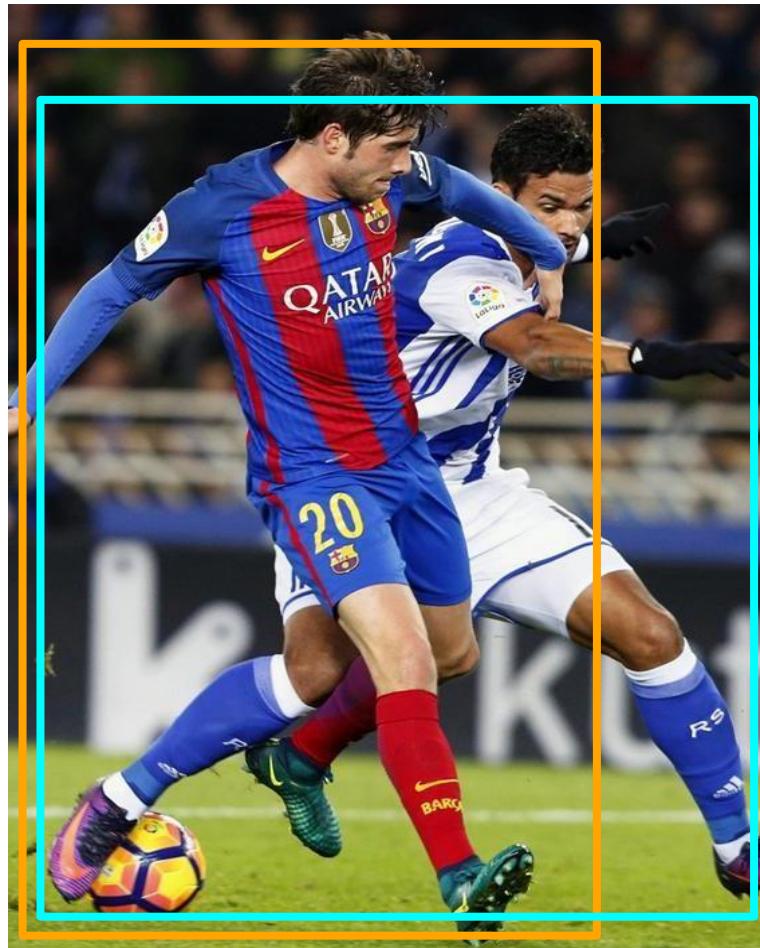




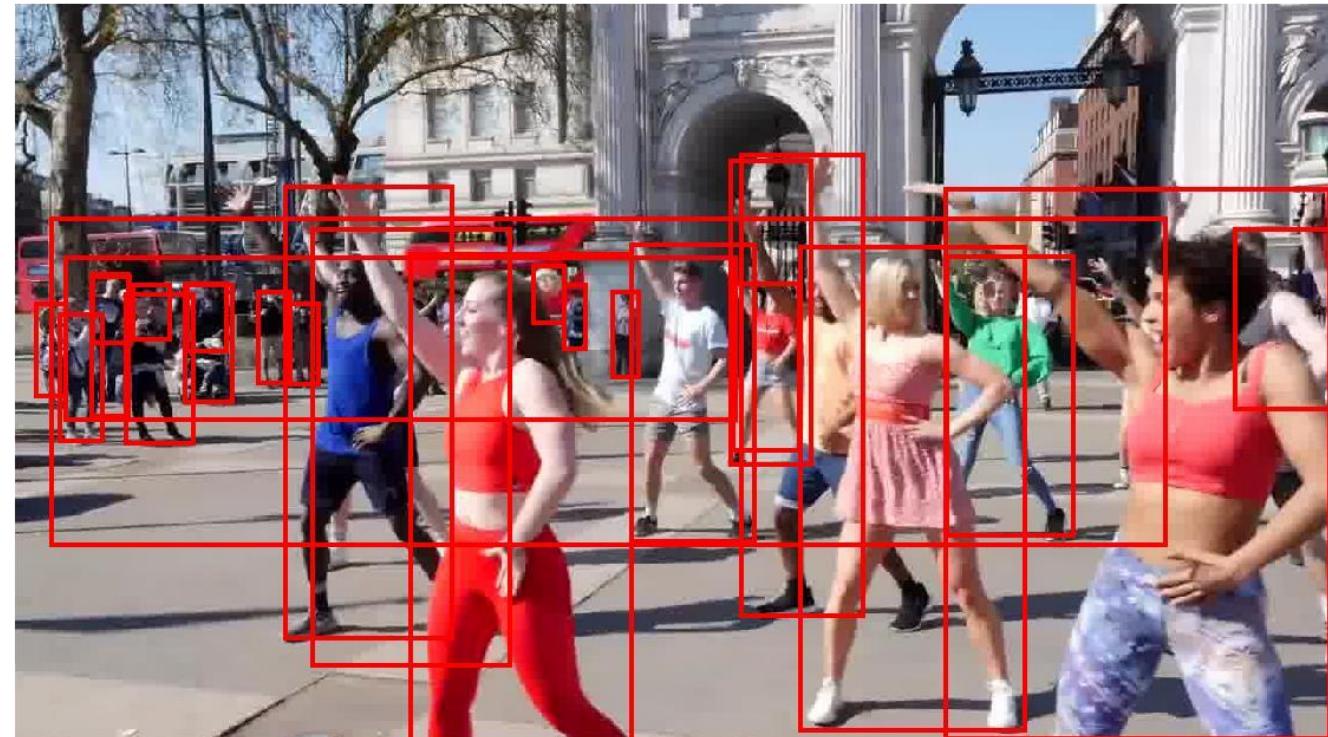
CrowdPose - Challenges



Input Human Proposals



- One Bounding Box contains **multi person**.
- Human Detection fails in crowded scenes.

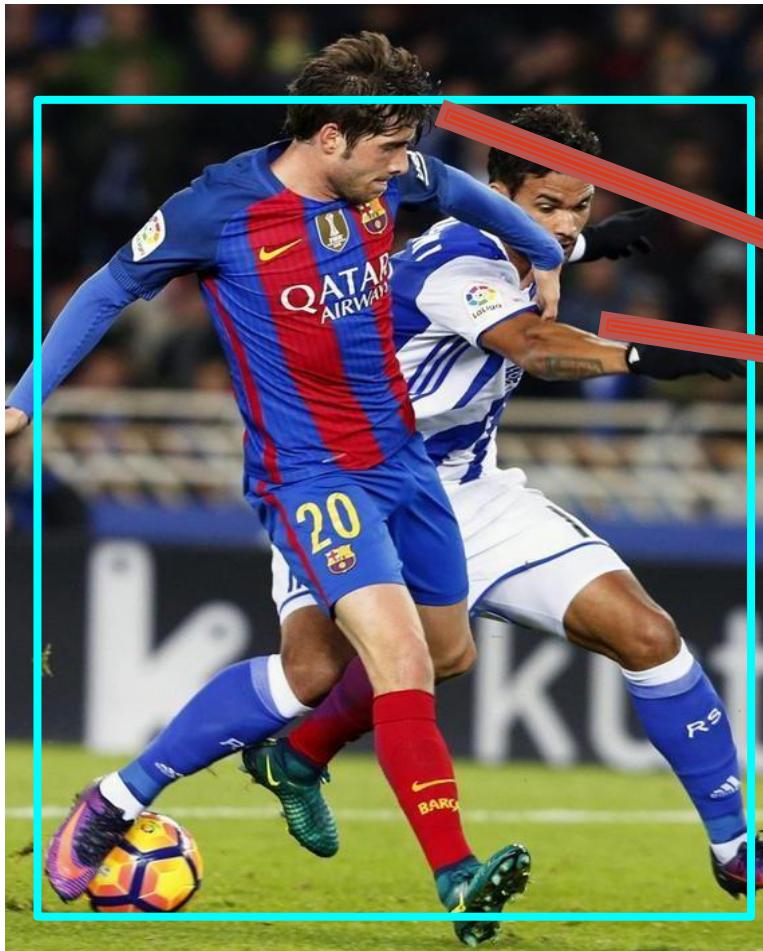


上海交通大学
SHANGHAI JIAO TONG UNIVERSITY

CrowdPose – Challenges



Input Human Proposals



- It is hard to identify which person this bounding box belongs to.

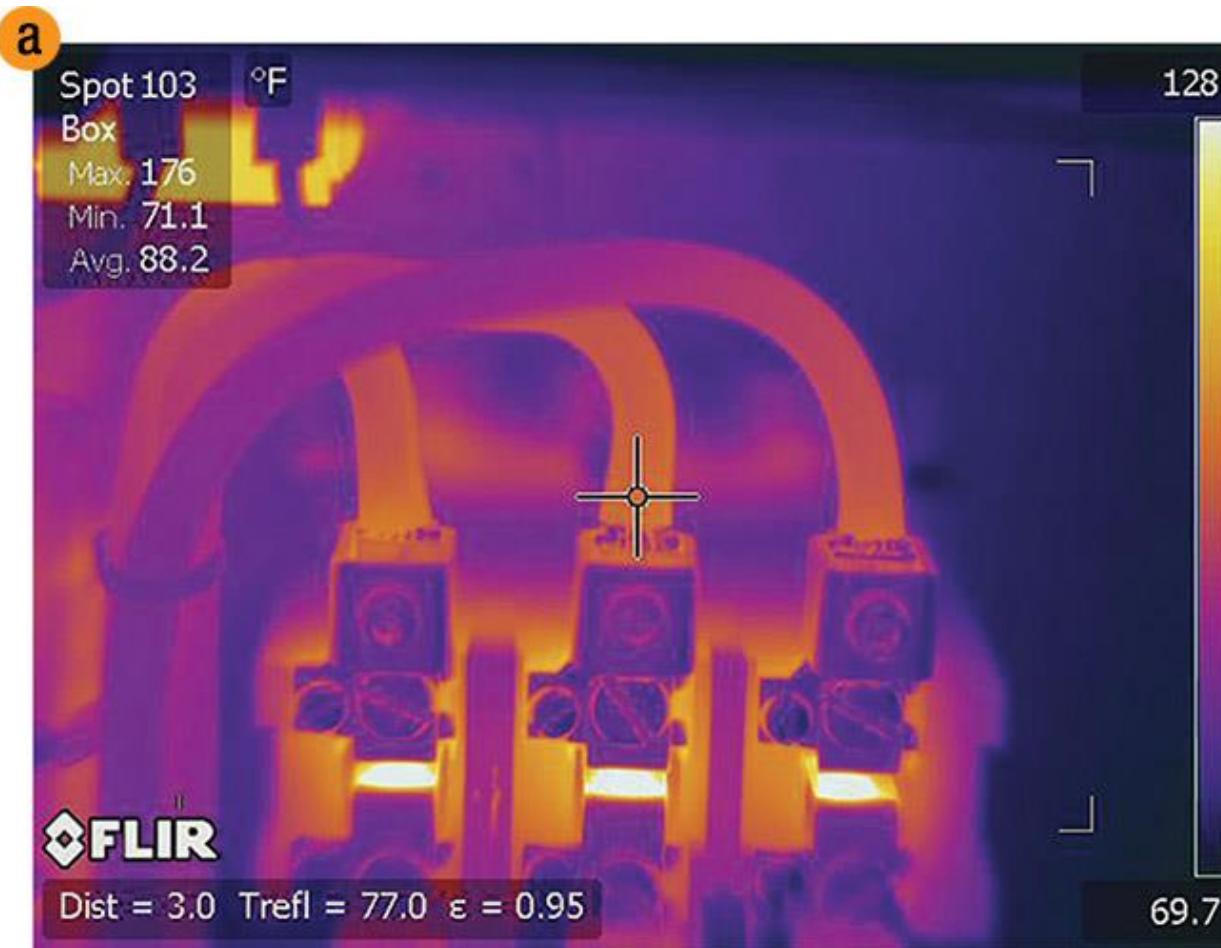
Person 1?

Person 2?



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY

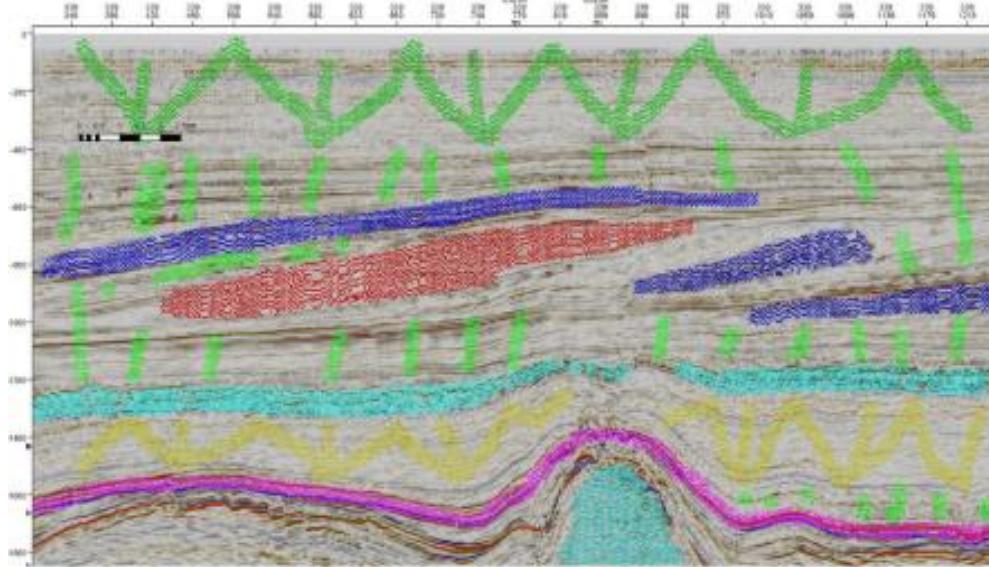
Infra-red image inspection



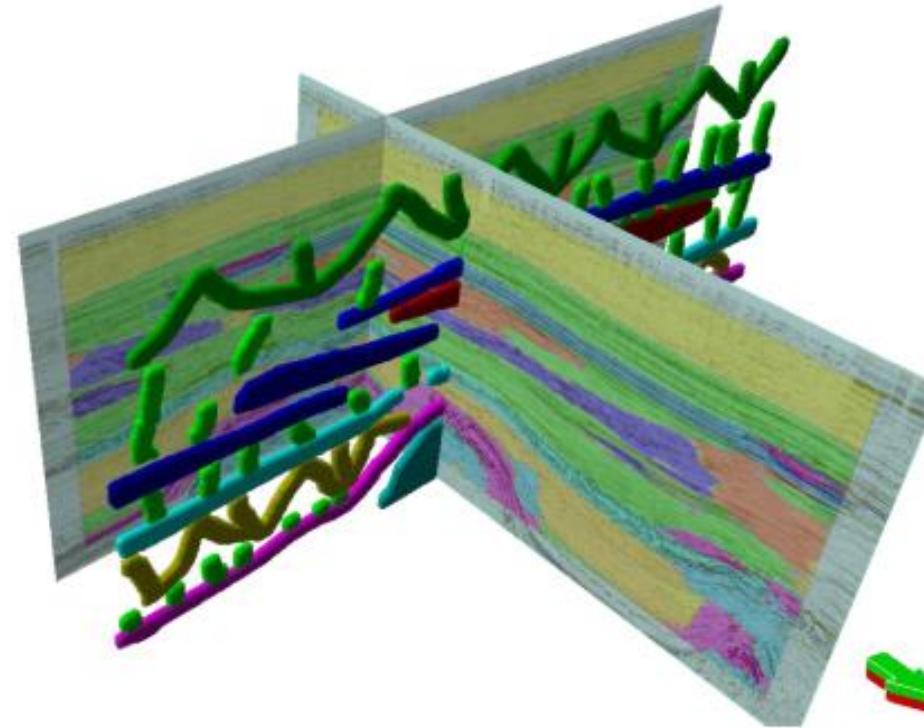
Seismic Data (3D)

Multi facies classification (F3 dataset)

One annotated inline with 9 facies classes



Auto-classified cube with 9 facies classes



Data provided by MalenoV project:

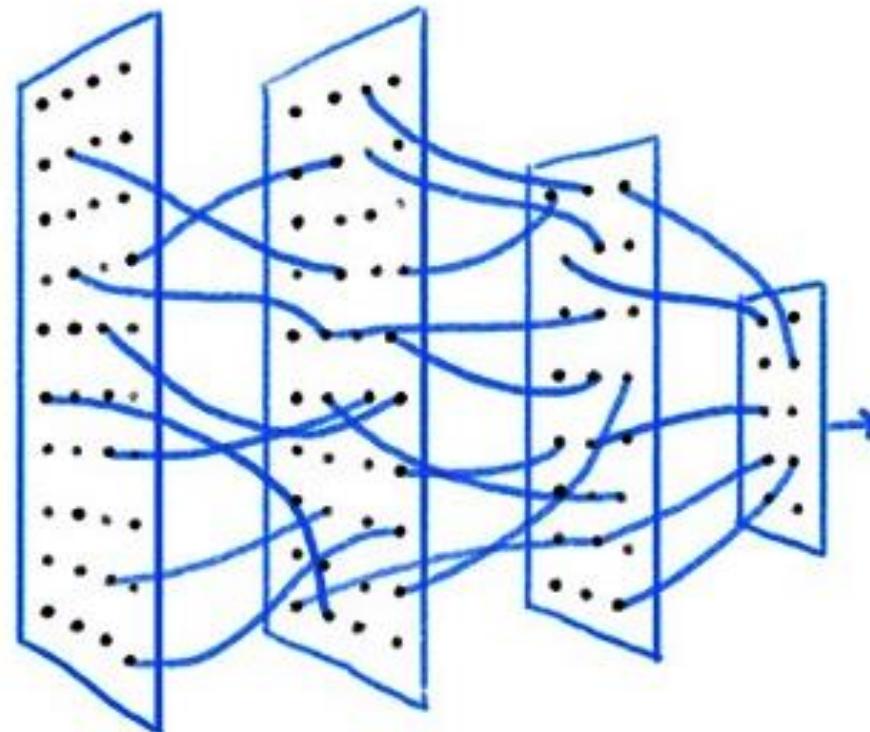
<https://github.com/bolgebrygg/MalenoV>

<https://github.com/microsoft/seismic-deeplearning>

Cat

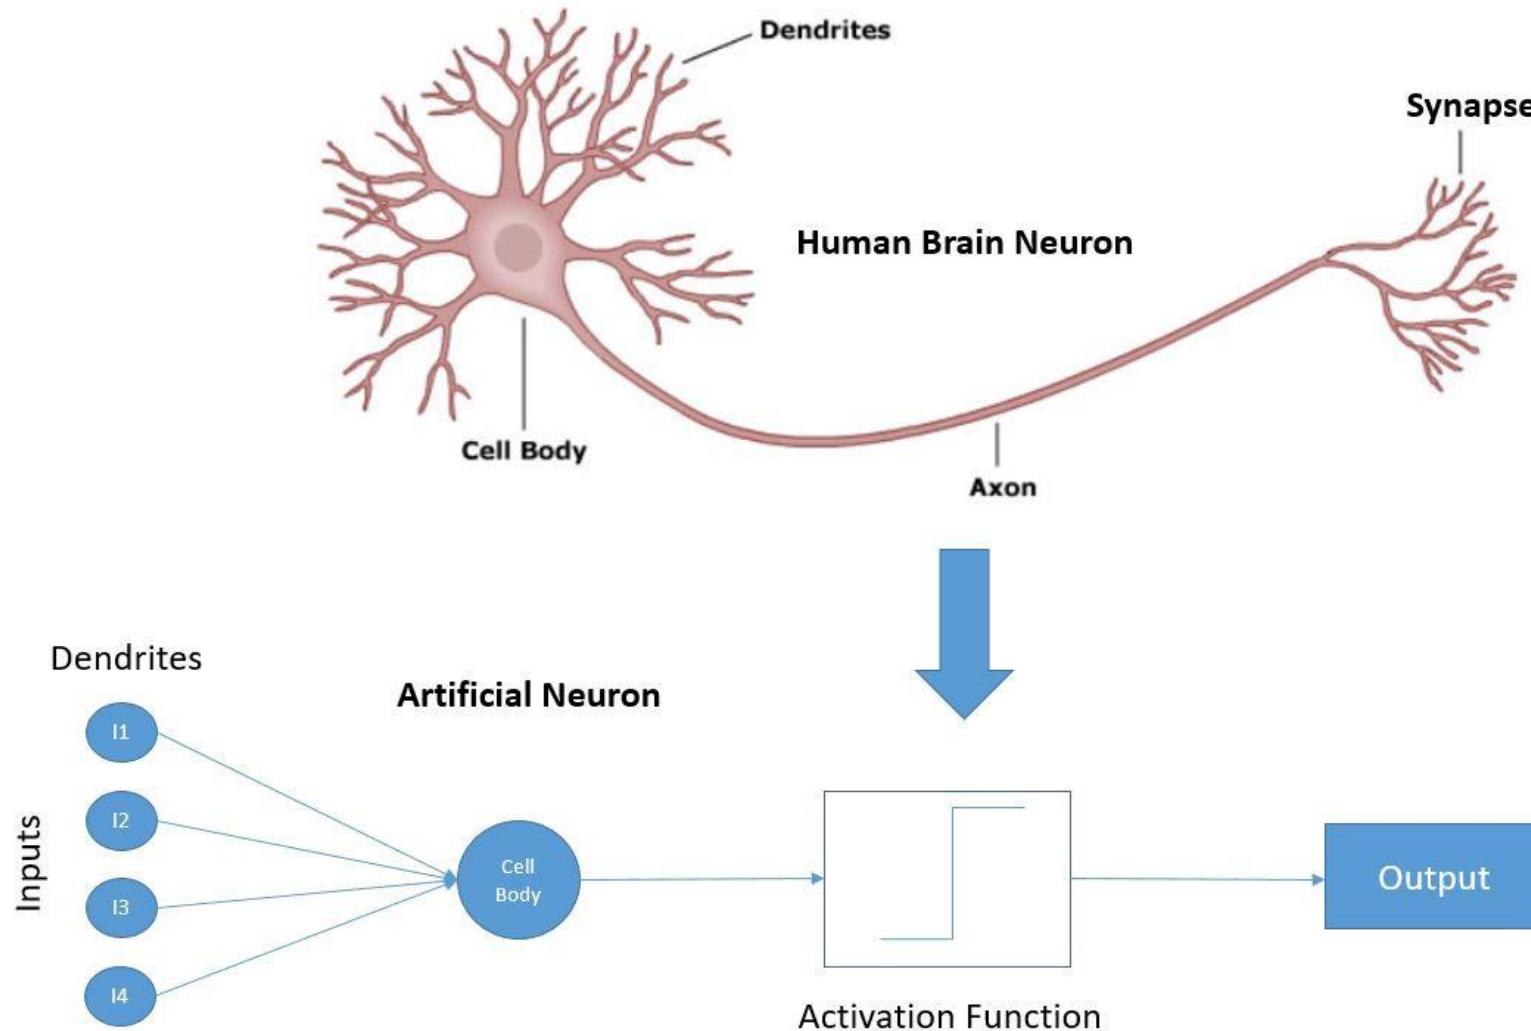


Dog

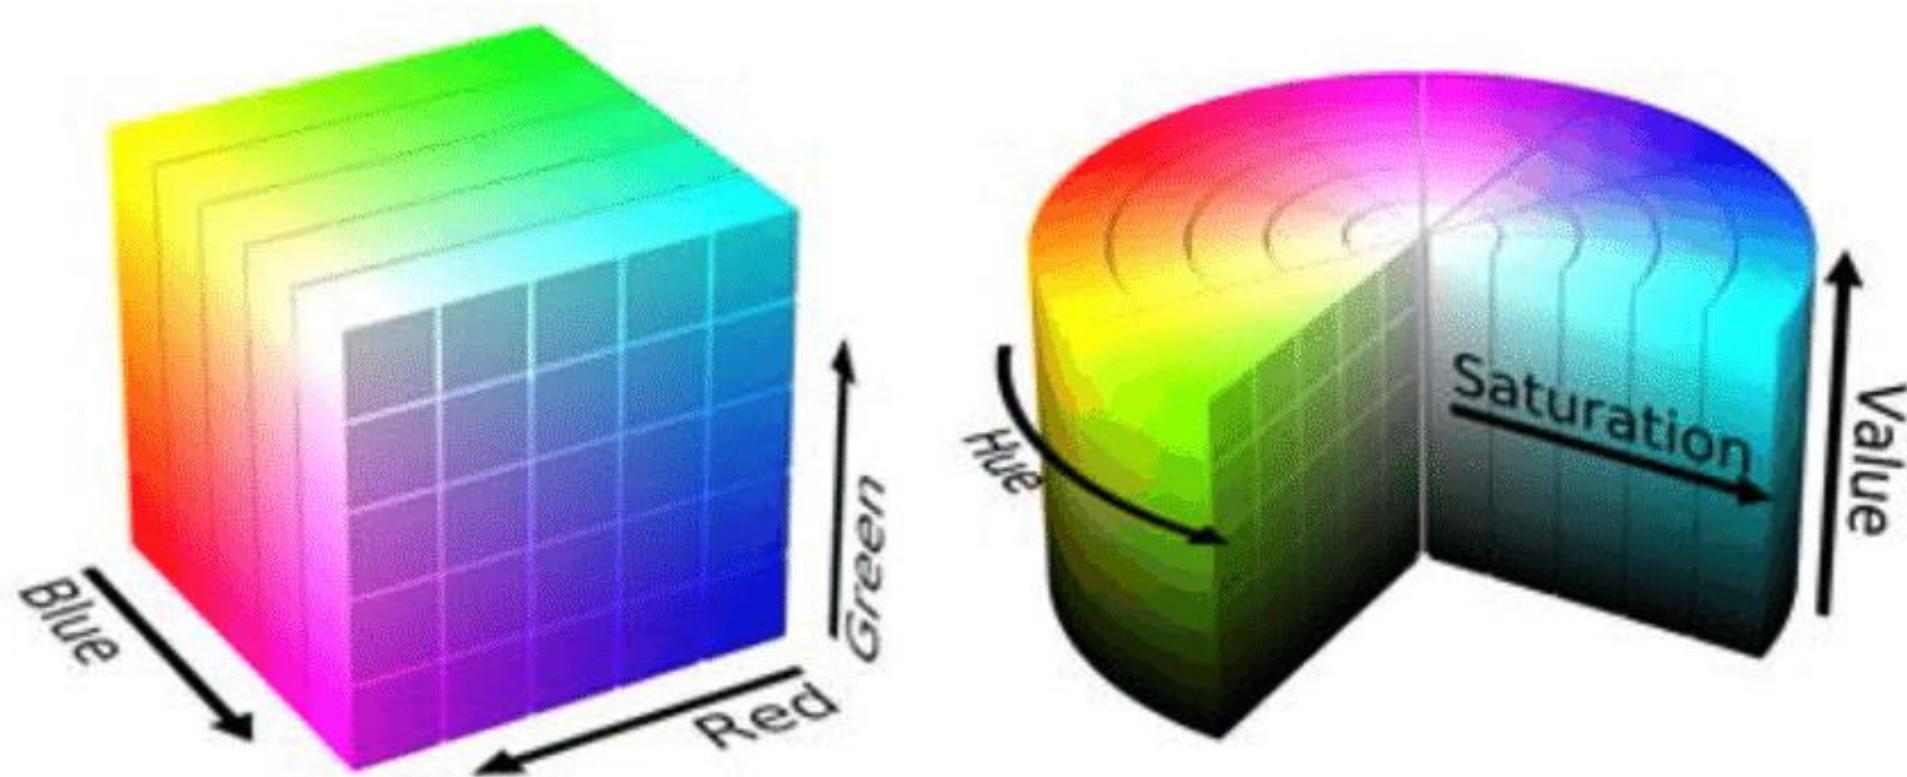


Output
Cat

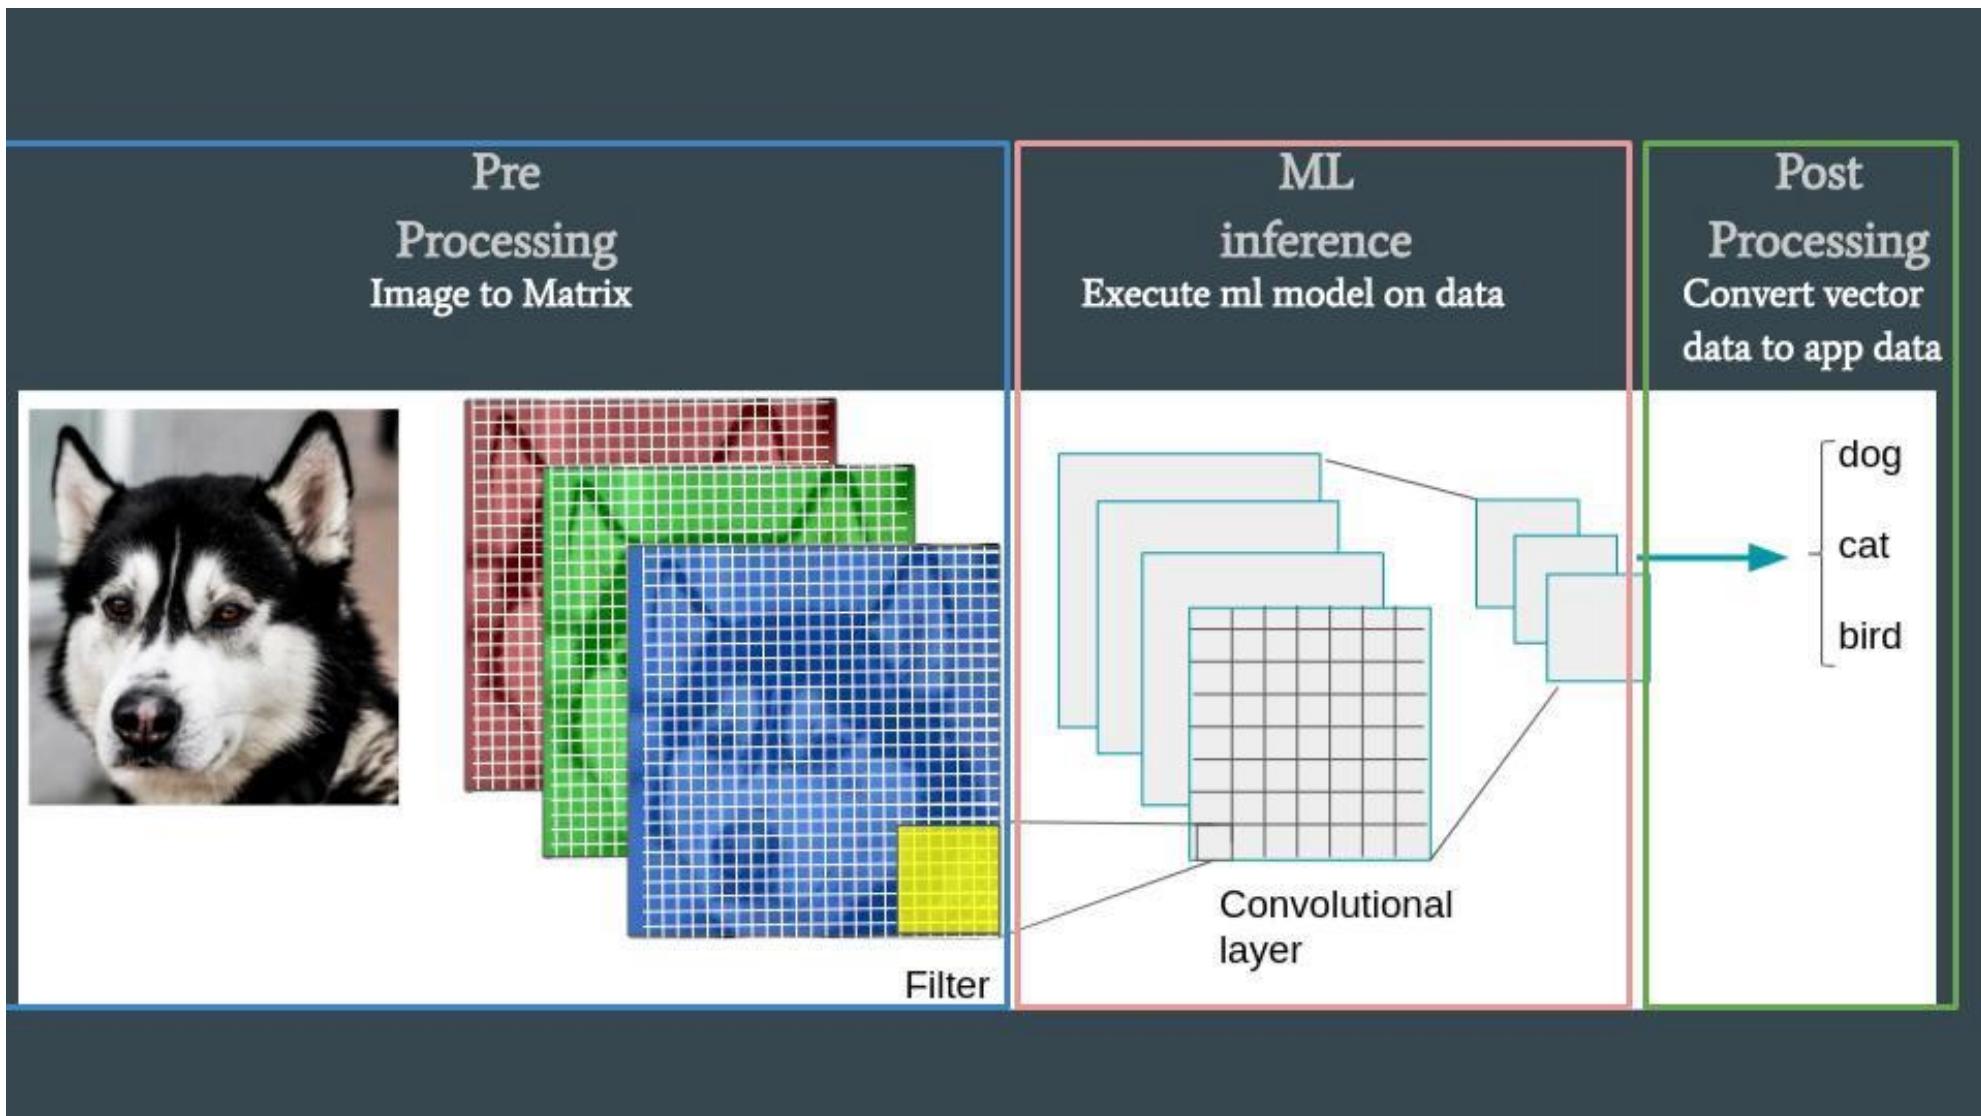
Why “neural” networks



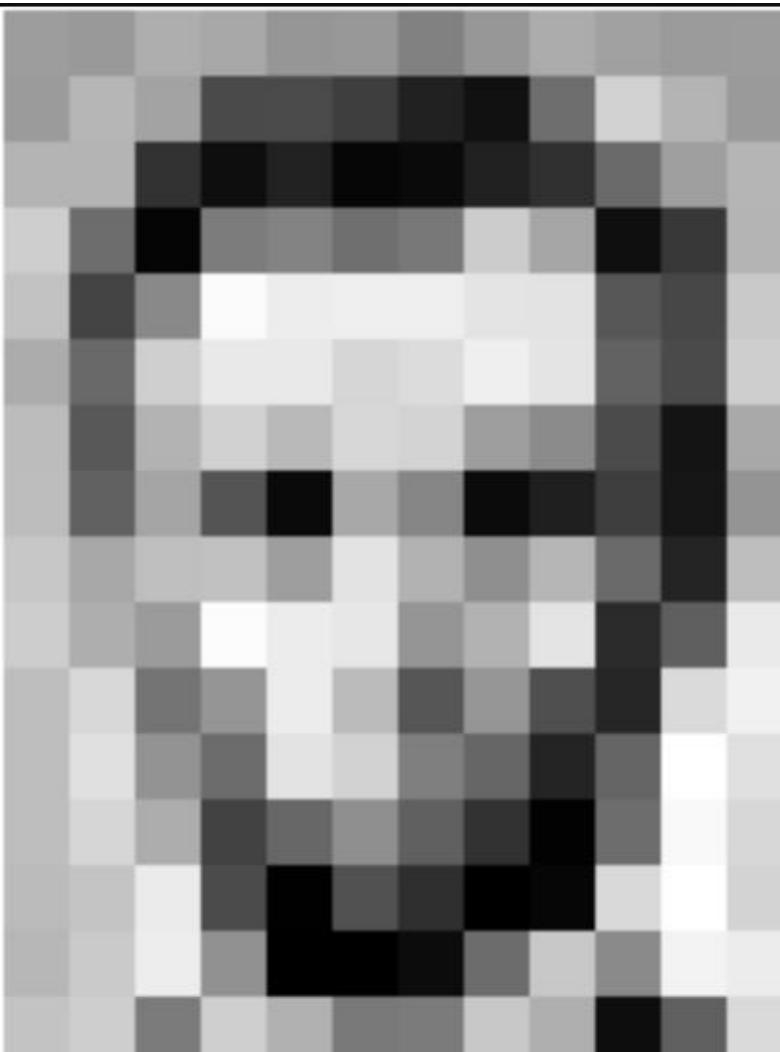
Color Space Representation



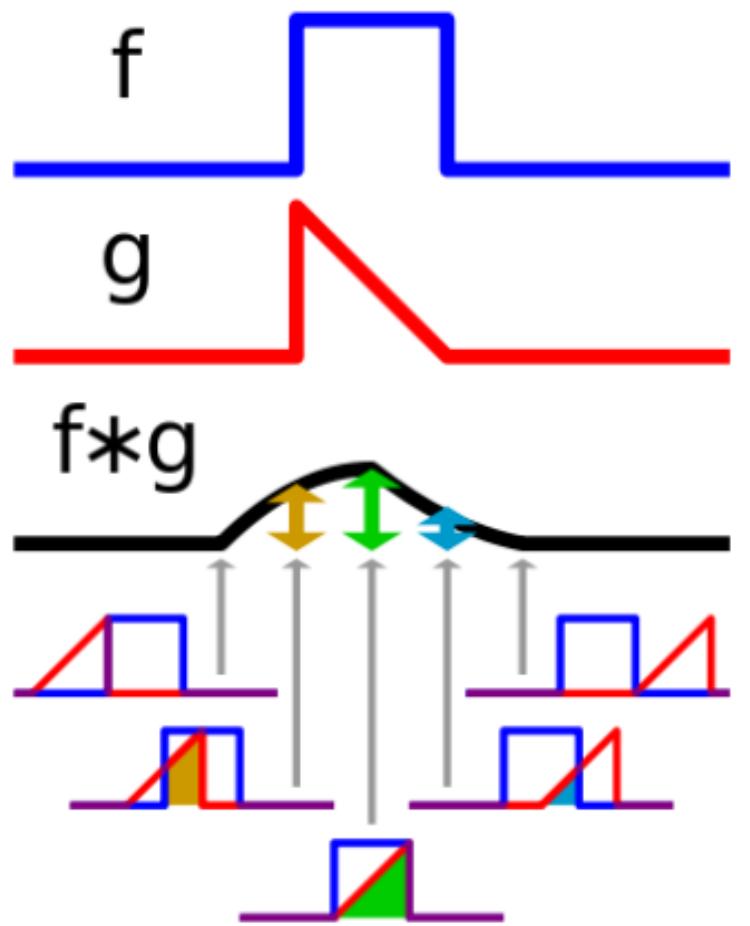
RGB image representation



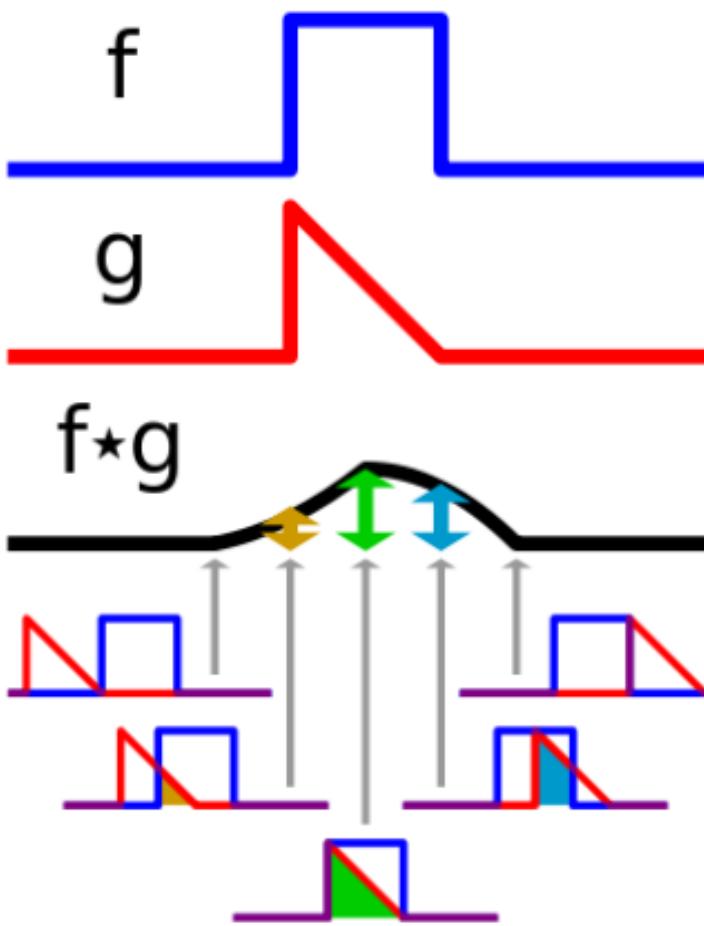
Grayscale Image Representation

																																																																																																																																																																																																																																																																																																																																																																																																	
<table border="1"><tr><td>157</td><td>153</td><td>174</td><td>168</td><td>150</td><td>152</td><td>129</td><td>151</td><td>172</td><td>161</td><td>155</td><td>156</td></tr><tr><td>155</td><td>182</td><td>163</td><td>74</td><td>75</td><td>62</td><td>33</td><td>17</td><td>110</td><td>210</td><td>180</td><td>154</td></tr><tr><td>180</td><td>180</td><td>50</td><td>14</td><td>34</td><td>6</td><td>10</td><td>33</td><td>48</td><td>105</td><td>159</td><td>181</td></tr><tr><td>206</td><td>109</td><td>5</td><td>124</td><td>131</td><td>111</td><td>120</td><td>204</td><td>166</td><td>15</td><td>56</td><td>180</td></tr><tr><td>194</td><td>68</td><td>137</td><td>251</td><td>237</td><td>239</td><td>239</td><td>228</td><td>227</td><td>87</td><td>71</td><td>201</td></tr><tr><td>172</td><td>106</td><td>207</td><td>233</td><td>233</td><td>214</td><td>220</td><td>239</td><td>228</td><td>98</td><td>74</td><td>206</td></tr><tr><td>188</td><td>88</td><td>179</td><td>209</td><td>185</td><td>215</td><td>211</td><td>158</td><td>199</td><td>75</td><td>20</td><td>169</td></tr><tr><td>189</td><td>97</td><td>165</td><td>84</td><td>10</td><td>168</td><td>134</td><td>11</td><td>31</td><td>62</td><td>22</td><td>148</td></tr><tr><td>199</td><td>168</td><td>191</td><td>193</td><td>158</td><td>227</td><td>178</td><td>143</td><td>182</td><td>105</td><td>36</td><td>190</td></tr><tr><td>206</td><td>174</td><td>155</td><td>252</td><td>236</td><td>231</td><td>149</td><td>178</td><td>228</td><td>43</td><td>95</td><td>234</td></tr><tr><td>190</td><td>216</td><td>116</td><td>149</td><td>236</td><td>187</td><td>85</td><td>150</td><td>79</td><td>38</td><td>218</td><td>241</td></tr><tr><td>190</td><td>224</td><td>147</td><td>108</td><td>227</td><td>210</td><td>127</td><td>102</td><td>35</td><td>101</td><td>255</td><td>224</td></tr><tr><td>190</td><td>214</td><td>173</td><td>66</td><td>103</td><td>143</td><td>95</td><td>50</td><td>2</td><td>109</td><td>249</td><td>215</td></tr><tr><td>187</td><td>196</td><td>235</td><td>75</td><td>1</td><td>81</td><td>47</td><td>0</td><td>6</td><td>217</td><td>255</td><td>211</td></tr><tr><td>183</td><td>202</td><td>237</td><td>145</td><td>0</td><td>0</td><td>12</td><td>108</td><td>200</td><td>138</td><td>243</td><td>236</td></tr><tr><td>195</td><td>206</td><td>123</td><td>207</td><td>177</td><td>121</td><td>123</td><td>200</td><td>175</td><td>13</td><td>96</td><td>218</td></tr></table>	157	153	174	168	150	152	129	151	172	161	155	156	155	182	163	74	75	62	33	17	110	210	180	154	180	180	50	14	34	6	10	33	48	105	159	181	206	109	5	124	131	111	120	204	166	15	56	180	194	68	137	251	237	239	239	228	227	87	71	201	172	106	207	233	233	214	220	239	228	98	74	206	188	88	179	209	185	215	211	158	199	75	20	169	189	97	165	84	10	168	134	11	31	62	22	148	199	168	191	193	158	227	178	143	182	105	36	190	206	174	155	252	236	231	149	178	228	43	95	234	190	216	116	149	236	187	85	150	79	38	218	241	190	224	147	108	227	210	127	102	35	101	255	224	190	214	173	66	103	143	95	50	2	109	249	215	187	196	235	75	1	81	47	0	6	217	255	211	183	202	237	145	0	0	12	108	200	138	243	236	195	206	123	207	177	121	123	200	175	13	96	218	<table border="1"><tr><td>157</td><td>153</td><td>174</td><td>168</td><td>150</td><td>152</td><td>129</td><td>151</td><td>172</td><td>161</td><td>155</td><td>156</td></tr><tr><td>155</td><td>182</td><td>163</td><td>74</td><td>75</td><td>62</td><td>33</td><td>17</td><td>110</td><td>210</td><td>180</td><td>154</td></tr><tr><td>180</td><td>180</td><td>50</td><td>14</td><td>34</td><td>6</td><td>10</td><td>33</td><td>48</td><td>105</td><td>159</td><td>181</td></tr><tr><td>206</td><td>109</td><td>5</td><td>124</td><td>131</td><td>111</td><td>120</td><td>204</td><td>166</td><td>15</td><td>56</td><td>180</td></tr><tr><td>194</td><td>68</td><td>137</td><td>251</td><td>237</td><td>239</td><td>239</td><td>228</td><td>227</td><td>87</td><td>71</td><td>201</td></tr><tr><td>172</td><td>106</td><td>207</td><td>233</td><td>233</td><td>214</td><td>220</td><td>239</td><td>228</td><td>98</td><td>74</td><td>206</td></tr><tr><td>188</td><td>88</td><td>179</td><td>209</td><td>185</td><td>215</td><td>211</td><td>158</td><td>199</td><td>75</td><td>20</td><td>169</td></tr><tr><td>189</td><td>97</td><td>165</td><td>84</td><td>10</td><td>168</td><td>134</td><td>11</td><td>31</td><td>62</td><td>22</td><td>148</td></tr><tr><td>199</td><td>168</td><td>191</td><td>193</td><td>158</td><td>227</td><td>178</td><td>143</td><td>182</td><td>105</td><td>36</td><td>190</td></tr><tr><td>206</td><td>174</td><td>155</td><td>252</td><td>236</td><td>231</td><td>149</td><td>178</td><td>228</td><td>43</td><td>95</td><td>234</td></tr><tr><td>190</td><td>216</td><td>116</td><td>149</td><td>236</td><td>187</td><td>85</td><td>150</td><td>79</td><td>38</td><td>218</td><td>241</td></tr><tr><td>190</td><td>224</td><td>147</td><td>108</td><td>227</td><td>210</td><td>127</td><td>102</td><td>35</td><td>101</td><td>255</td><td>224</td></tr><tr><td>190</td><td>214</td><td>173</td><td>66</td><td>103</td><td>143</td><td>95</td><td>50</td><td>2</td><td>109</td><td>249</td><td>215</td></tr><tr><td>187</td><td>196</td><td>235</td><td>75</td><td>1</td><td>81</td><td>47</td><td>0</td><td>6</td><td>217</td><td>255</td><td>211</td></tr><tr><td>183</td><td>202</td><td>237</td><td>145</td><td>0</td><td>0</td><td>12</td><td>108</td><td>200</td><td>138</td><td>243</td><td>236</td></tr><tr><td>195</td><td>206</td><td>123</td><td>207</td><td>177</td><td>121</td><td>123</td><td>200</td><td>175</td><td>13</td><td>96</td><td>218</td></tr></table>	157	153	174	168	150	152	129	151	172	161	155	156	155	182	163	74	75	62	33	17	110	210	180	154	180	180	50	14	34	6	10	33	48	105	159	181	206	109	5	124	131	111	120	204	166	15	56	180	194	68	137	251	237	239	239	228	227	87	71	201	172	106	207	233	233	214	220	239	228	98	74	206	188	88	179	209	185	215	211	158	199	75	20	169	189	97	165	84	10	168	134	11	31	62	22	148	199	168	191	193	158	227	178	143	182	105	36	190	206	174	155	252	236	231	149	178	228	43	95	234	190	216	116	149	236	187	85	150	79	38	218	241	190	224	147	108	227	210	127	102	35	101	255	224	190	214	173	66	103	143	95	50	2	109	249	215	187	196	235	75	1	81	47	0	6	217	255	211	183	202	237	145	0	0	12	108	200	138	243	236	195	206	123	207	177	121	123	200	175	13	96	218
157	153	174	168	150	152	129	151	172	161	155	156																																																																																																																																																																																																																																																																																																																																																																																						
155	182	163	74	75	62	33	17	110	210	180	154																																																																																																																																																																																																																																																																																																																																																																																						
180	180	50	14	34	6	10	33	48	105	159	181																																																																																																																																																																																																																																																																																																																																																																																						
206	109	5	124	131	111	120	204	166	15	56	180																																																																																																																																																																																																																																																																																																																																																																																						
194	68	137	251	237	239	239	228	227	87	71	201																																																																																																																																																																																																																																																																																																																																																																																						
172	106	207	233	233	214	220	239	228	98	74	206																																																																																																																																																																																																																																																																																																																																																																																						
188	88	179	209	185	215	211	158	199	75	20	169																																																																																																																																																																																																																																																																																																																																																																																						
189	97	165	84	10	168	134	11	31	62	22	148																																																																																																																																																																																																																																																																																																																																																																																						
199	168	191	193	158	227	178	143	182	105	36	190																																																																																																																																																																																																																																																																																																																																																																																						
206	174	155	252	236	231	149	178	228	43	95	234																																																																																																																																																																																																																																																																																																																																																																																						
190	216	116	149	236	187	85	150	79	38	218	241																																																																																																																																																																																																																																																																																																																																																																																						
190	224	147	108	227	210	127	102	35	101	255	224																																																																																																																																																																																																																																																																																																																																																																																						
190	214	173	66	103	143	95	50	2	109	249	215																																																																																																																																																																																																																																																																																																																																																																																						
187	196	235	75	1	81	47	0	6	217	255	211																																																																																																																																																																																																																																																																																																																																																																																						
183	202	237	145	0	0	12	108	200	138	243	236																																																																																																																																																																																																																																																																																																																																																																																						
195	206	123	207	177	121	123	200	175	13	96	218																																																																																																																																																																																																																																																																																																																																																																																						
157	153	174	168	150	152	129	151	172	161	155	156																																																																																																																																																																																																																																																																																																																																																																																						
155	182	163	74	75	62	33	17	110	210	180	154																																																																																																																																																																																																																																																																																																																																																																																						
180	180	50	14	34	6	10	33	48	105	159	181																																																																																																																																																																																																																																																																																																																																																																																						
206	109	5	124	131	111	120	204	166	15	56	180																																																																																																																																																																																																																																																																																																																																																																																						
194	68	137	251	237	239	239	228	227	87	71	201																																																																																																																																																																																																																																																																																																																																																																																						
172	106	207	233	233	214	220	239	228	98	74	206																																																																																																																																																																																																																																																																																																																																																																																						
188	88	179	209	185	215	211	158	199	75	20	169																																																																																																																																																																																																																																																																																																																																																																																						
189	97	165	84	10	168	134	11	31	62	22	148																																																																																																																																																																																																																																																																																																																																																																																						
199	168	191	193	158	227	178	143	182	105	36	190																																																																																																																																																																																																																																																																																																																																																																																						
206	174	155	252	236	231	149	178	228	43	95	234																																																																																																																																																																																																																																																																																																																																																																																						
190	216	116	149	236	187	85	150	79	38	218	241																																																																																																																																																																																																																																																																																																																																																																																						
190	224	147	108	227	210	127	102	35	101	255	224																																																																																																																																																																																																																																																																																																																																																																																						
190	214	173	66	103	143	95	50	2	109	249	215																																																																																																																																																																																																																																																																																																																																																																																						
187	196	235	75	1	81	47	0	6	217	255	211																																																																																																																																																																																																																																																																																																																																																																																						
183	202	237	145	0	0	12	108	200	138	243	236																																																																																																																																																																																																																																																																																																																																																																																						
195	206	123	207	177	121	123	200	175	13	96	218																																																																																																																																																																																																																																																																																																																																																																																						

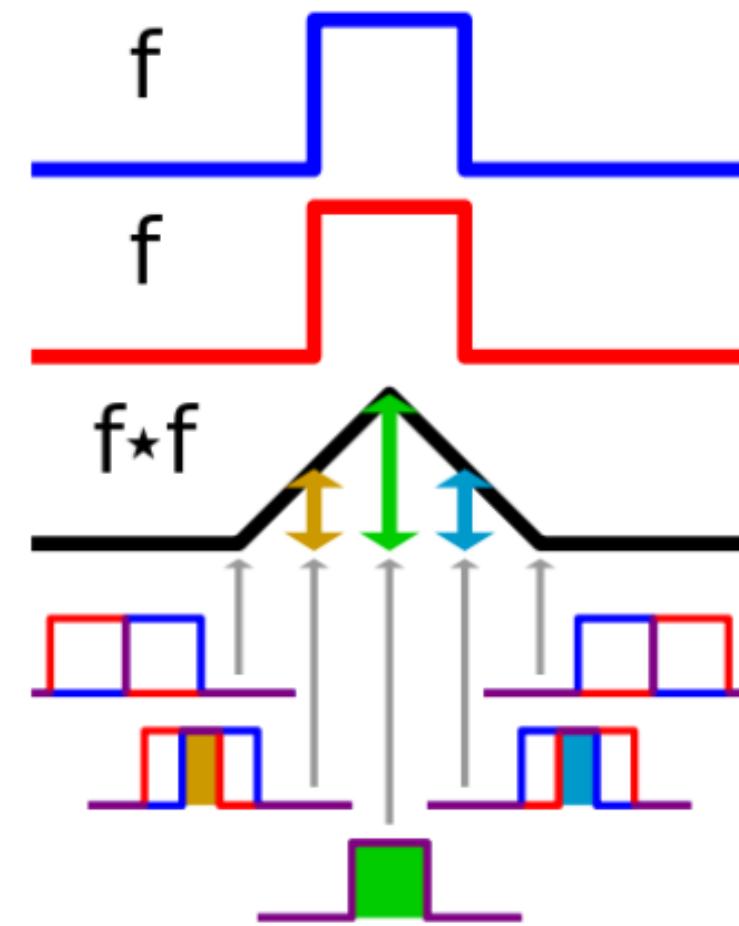
Convolution



Cross-correlation



Autocorrelation



Convolution and Cross-Correlation will just flip the weight kernel for NNs!!

CNN recap

$$y_i = \frac{e^{x_i}}{\sum_k e^{x_k}}$$

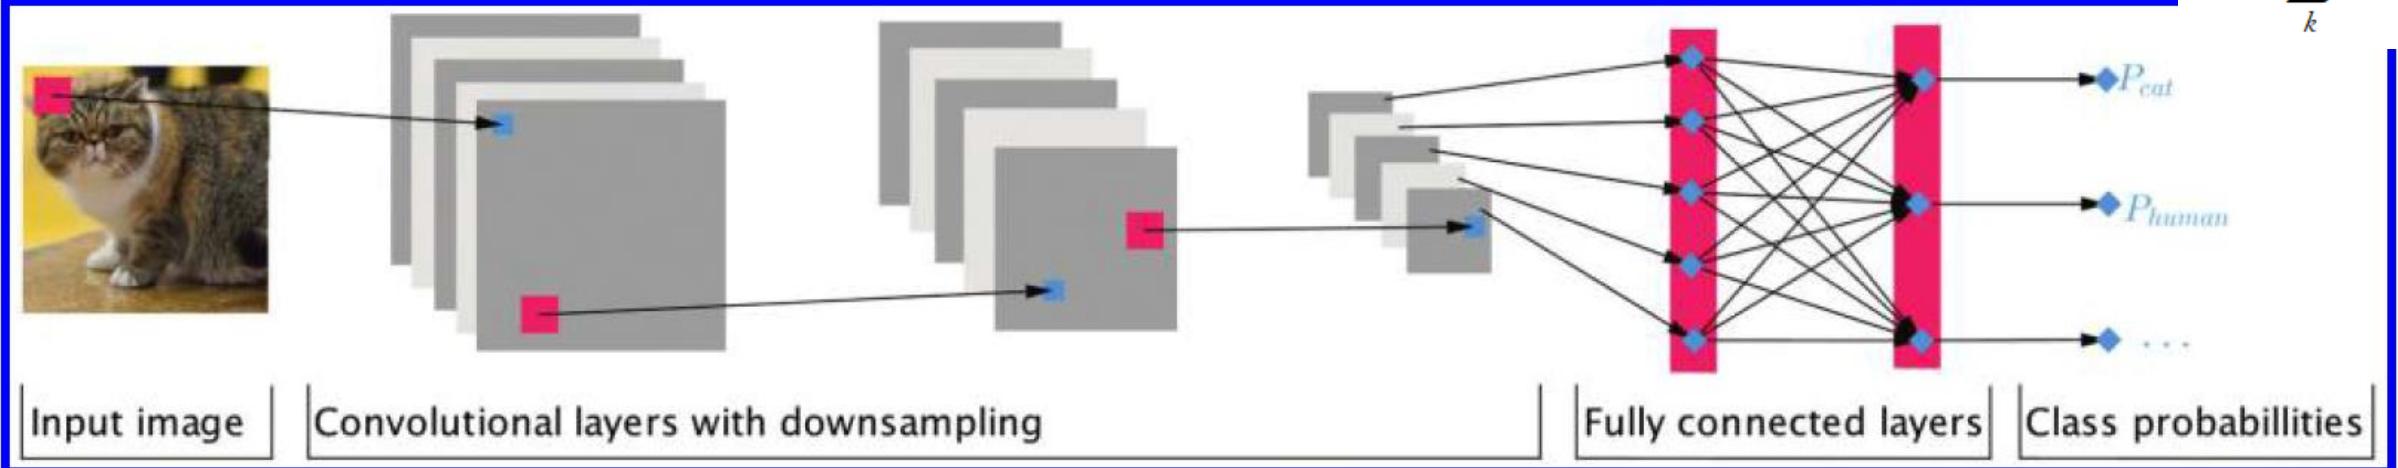


Figure 3. A CNN applied for image classification. The input is an image, and the output is a probability vector in which each element contains the estimated probability of the image containing an object of the given class.

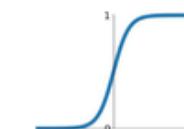
$$X_{i+1}(x, y, j) = f \left(\sum_{k=1}^{C_{i-1}} X_i(x, y, k) * W_{i,j}(x, y, k) \right)$$

Activation function f can be linear or non-linear
 C is the channel dimension
Zero padding is used to resize image to common dimension

Activation Functions

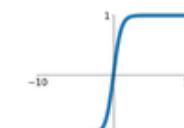
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



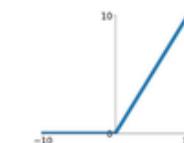
tanh

$$\tanh(x)$$



ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$



Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

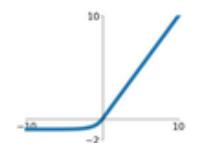


Image Classification demo

https://github.com/microsoft/computervision-recipes/blob/master/scenarios/classification/01_training_introduction.ipynb

- More CV scenarios can be found at
- <https://github.com/microsoft/computervision-recipes/tree/master/scenarios>

References

- <https://www.themtank.org/a-year-in-computer-vision>
- <https://www.analyticsvidhya.com/blog/2020/09/18-open-source-computer-vision-projects-beginners/>
- <https://www.kaggle.com/hrmello/intro-to-image-processing-colorspaces>
- <https://becominghuman.ai/real-computer-vision-for-mobile-and-embedded-part-3-postprocessing-96966f707fbd?gi=375e9364dfe8>